

Inside Goldman's Oracle ATS

The complete machine: why Goldman runs on Oracle, how Oracle Recruiting Cloud processes you step by step, and what "scoring highly" actually means

Companion to "The London Map & The Real ATS" · badges: **VERIFIED** = Oracle docs / GS's own system / on-record reporting · **LIKELY** = credible secondary or standard product behaviour · **BEST GUESS** = labelled inference about GS's private configuration

The one reframe that makes everything below useful

There is no single "ATS score" waiting to be gamed. Oracle Recruiting Cloud is not a judge — it is a *pipeline*: identity → parsing → knockout questions → a staged selection workflow → recruiter tools. Each stage can only do one of three things to you: **block you** (mechanically), **garble you** (parsing), or **surface you** (search & matching).

So "scoring highly" decomposes into four real, controllable objectives: **never trip a knockout · parse perfectly · be findable · read well in the human seconds that follow.** Everything in this document serves one of those four.

1. Why Goldman runs on Oracle — the story almost nobody knows

This is not a generic "banks like Oracle" story. The documented sequence: **VERIFIED (on-record reporting)**

1. **Goldman first bought Workday** — the fashionable cloud HR system most of its peers use. The implementation **did not succeed**.
2. **Goldman then switched to Oracle Fusion Cloud HCM**, going live around 2020 — a defection Oracle's leadership publicly celebrated at Workday's expense, which is why it's on the record at all.
3. In the same wave, Goldman **selected Oracle Cloud HCM Recruiting (ORC) as its ATS**, displacing legacy recruiting systems — the banking-era standard it replaced was Taleo-class technology **LIKELY**.

Why Oracle rather than anything else? **BEST GUESS — reasoned, not confirmed**

- **One suite, one record.** Fusion HCM runs core HR (payroll, workforce, talent) and recruiting in a single system — a candidate who becomes an employee never changes databases. At 40k+ employees and ~500k applications a year across programmes, that continuity is worth more than a prettier ATS.
- **Skinnability.** ORC's CandidateExperience layer is built to be white-labelled — which is exactly what `higher.gs.com` is. Goldman keeps its brand; Oracle keeps the plumbing.
- **Configurable process, not fixed product.** ORC's selection workflows, question banks and automation rules are tenant-configured — Goldman can encode its own rules (the 4-application cap behaviour, campus vs lateral sites) without building software.
- **Bank-grade compliance.** Auditability of every candidate decision, data residency, regulator-friendly reporting — Oracle's home turf.

Why you should care: because it means the system screening you is a *configured product with documented behaviour*, not a mystery black box. Oracle publishes exactly what ORC can and cannot do — that documentation is what the rest of this PDF is built on. The only unknown is which optional features Goldman has switched on, and I flag that wherever it matters.

2. The architecture — what talks to what

WHAT YOU SEE	higher.gs.com (branded career site) → "CandidateExperience" pages on hdpc.fa.us2.oraclecloud.com Search, role pages, the apply flow. The /sites/LateralHiring/ vs campus distinction is cosmetic skin routing — same engine behind both. VERIFIED
IDENTITY	Email-first candidate account (/apply/email) Email → verification code → profile created or matched. Your email is the primary key for everything: applications, HireVue invites, duplicate detection. VERIFIED
THE ENGINE	Oracle Recruiting Cloud core (Goldman tenant "hdpc", Oracle US2 region) Requisitions, applications, parsed profiles, question banks, the candidate selection workflow, recruiter tooling, all reporting. VERIFIED
BOLT-ONS	Assessment & interview partners wired into the workflow HireVue (video interview), HackerRank (Engineering) — triggered as workflow stages, delivered by email. ORC supports auto-progression on assessment results (see §4). VERIFIED that these are used; wiring detail LIKELY
HUMANS	Recruiting teams working per-requisition queues Rolling review of the pipeline per role — the "different roles at different times" behaviour you see from outside. VERIFIED

3. Your application's journey through the machine, stage by stage

1	Account & identity Email verification creates/matches your candidate record. Duplicate detection is native — same person, multiple emails is exactly what it's built to catch, and Goldman's posting explicitly forbids it. VERIFIED
2	Resume parsing → structured profile On upload, ORC's parser extracts education, employers, titles, dates and skills into database fields — this parsed profile, not your pretty PDF, is what search, filters and any matching features operate on. The PDF itself stays attached for human eyes. ORC shows you the extracted fields during apply: correcting any mis-parse by hand is a scoring action almost nobody bothers with. VERIFIED (core ORC feature)
3	Prescreening questionnaire — including disqualification questions ORC lets employers attach question sets per requisition. A disqualification question is a single-answer question encoding a minimum requirement — "answers decide if candidates move forward or are automatically disqualified" (Oracle's own wording). For campus: graduation window, penultimate/final-year status, work authorisation. This is the only true auto-reject in the system. VERIFIED (ORC feature) · GS's exact question set BEST GUESS
4	Entry into the Candidate Selection Process (CSP) Every requisition runs a configured workflow of phases, each with states — canonically New → Screening → Interview(s) → Offer → HR. Your application sits at exactly one phase/state at all times; recruiters move you, or automation does. VERIFIED (ORC architecture)

Automated progression rules

- 5 ORC supports auto-advance/auto-hold conditions (Oracle "fast formulas") — e.g. "candidate completes and passes every assessment in a series → move to next phase." This is how HireVue completion can trigger the next stage without a human touching it. Whether GS automates progression or keeps humans on every move: unknown. **VERIFIED capability** · GS usage **UNKNOWN**

Recruiter workspace — where the real triage happens

- 6 Recruiters see the requisition pipeline with parsed profiles, filters, keyword search, prescreen answers, and (if licensed) AI features. This is the human seconds-long read every previous document described — now you know the screen it happens on. **VERIFIED tooling**

Stage invitations pushed by email

- 7 HireVue link, interview scheduling, offer — all keyed to the email from step 1. Short completion windows are configured per stage. **VERIFIED**

New
application lands

Screening
prescreen + triage

Assessment
HireVue /
HackerRank

Interview
Superday rounds

Offer
drafted & approved

HR
onboarding

The canonical ORC phase sequence mapped onto Goldman's known campus funnel. Names/states inside GS's tenant are their own configuration **BEST GUESS**

4. Does the machine rank you? The honest anatomy

Four distinct mechanisms get conflated as "the ATS ranks your CV." Separate them and the fog clears:

Mechanism	Exists in ORC?	On at Goldman campus?	What it actually does
Disqualification questions	YES	Almost certainly	Binary knockout on eligibility answers. Not a ranking — a gate. The only automated rejection in the product.
Keyword search & filters (recruiter-side)	YES	Standard tooling	Recruiters query the parsed profiles ("Python", "M&A", school names). Being <i>findable</i> = conventional vocabulary in parseable text. Absence from a search isn't rejection — it's invisibility.
AI matching / compatibility ratings ("Recruiting Booster": Suggested Candidates, AI fit indicators)	YES, as paid add-on	UNKNOWN	Where licensed, compares parsed profile against requisition text and surfaces suggested/rated candidates <i>to the recruiter</i> — it assists shortlisting, it does not auto-reject. Whether GS runs it for campus is not public. Note GS's own posting text is thin (empty qualification fields), which limits what any matcher could match against — a real structural point in your favour.
A hidden quality score that auto-rejects CVs	NO	—	Not how the product works. Quality judgement in this pipeline is human triage in the recruiter workspace. Every earlier conclusion stands.

The myth, killed precisely: when someone says "the ATS rejected my CV," what actually happened in an ORC shop is almost always one of: they tripped a disqualification question; their CV parsed badly so no search ever surfaced them; or a human spent eight seconds and moved on. Three different failures, three different fixes — and none of them is "add more keywords in white text" (parsed text is what recruiters see; invisible-text tricks are visible in the parsed profile and read as cheating).

5. Scoring highly — the playbook, mapped to the machine

Objective 1 — Never trip a knockout highest stakes, fully controllable

- Answer eligibility questions **exactly accurately** — graduation date, year of study, right to work. A careless dropdown answer is the one genuine auto-reject available.
- Know your graduation-year maths before you start (2028 grads = penultimate year = this cycle).
- Stay inside the **4-application cap**; a 5th auto-withdraws an earlier one. One email address, forever.
- If a question is ambiguous for your situation (placement year, integrated Master's), resolve it before submitting — a wrong guess here ends the application unread.

Objective 2 — Parse perfectly the mechanical layer

- One page, single column, no tables/text boxes/graphics, standard headings, text-selectable PDF (run the select-and-copy test).
- Dates in one consistent format on every entry — parsers build your work history from them; inconsistency creates gaps in the structured record.
- Company name and job title on separate, plain lines (your Tracker template already does this).
- **The unglamorous edge: when ORC displays parsed fields during apply, proof-read and fix them.** Your searchable record is only as good as what landed in the database.
- Degree, university and graduation year written plainly — these are the most-queried fields in the system.

Objective 3 — Be findable the searchability layer

- Use **conventional nouns** for institutions, roles and skills: "Client Onboarding Intern, Pictet" beats a creative title. Recruiters search standard vocabulary.
- Name real, hard skills and tools explicitly (Excel, Python, Bloomberg, languages) — these are filter terms.
- Mirror the requisition's genuine nouns where truthful (for CSG: client onboarding, institutional clients, pension funds, mandates) — this serves human fluency AND any matching algorithm, since both compare your parsed text to the requisition's text. Same rule as ever: ideas in their vocabulary, never pasted phrases, nothing you can't defend.
- No keyword stuffing, no white text — the parsed profile is literally what the recruiter reads in the workspace. Stuffing is self-reporting that you tried to cheat.

Objective 4 — Win the human seconds where the actual decision happens

- Everything from your CV masterclass applies at the moment a recruiter opens your record: strongest client-facing bullet first, quantified facts, division alignment visible in five seconds.
- Prescreen answers, CV and motivation statement must tell **one consistent story** — the recruiter sees them side by side on one screen, and downstream interviewers see all of it against your HireVue.
- The motivation statement is the only free-text the triage read lingers on — that's why the trust/conductor spine you've built is the highest-leverage asset you have inside this entire machine.

6. Myths vs machine — quick-reference table

The myth	What the machine actually does
"The ATS scores every CV and rejects the bottom 90%."	Only disqualification questions auto-reject. CV quality is judged by humans in the recruiter workspace.
"You must copy the job description's keywords or a robot bins you."	GS campus postings have near-empty qualification fields — there's barely anything to match. Keywords matter for recruiter <i>search</i> and human fluency, not survival.

"Applying in the first hour boosts your algorithm rank."	No such rank exists. Timing matters only through rolling human capacity (covered fully in the mechanics PDF).
"PDFs break the ATS — use Word."	ORC parses text-based PDFs fine. What breaks parsing is layout (columns, tables, images), not the extension.
"A second account gives you more shots."	Duplicate detection is native, the posting bans it, and it reads as dishonesty on a permanent record.
"Once submitted, it's a black box."	It's a phase/state pipeline. Silence usually means "sitting in Screening at a rolling queue," not rejection. Watch the email you applied with.

7. The 60-second version

Goldman's ATS is Oracle Recruiting Cloud behind a Goldman skin — chosen after a failed Workday implementation, because one suite runs all of HR. It parses your CV into a database record, asks knockout eligibility questions (the only auto-reject), then routes you through a phase-by-phase workflow where humans triage per role on a rolling basis, assisted by search and possibly AI suggestions. You "score highly" by: answering knockouts exactly right, parsing cleanly and fixing the parsed fields, writing in conventional searchable vocabulary that mirrors the desk's real language, and making the human seconds count with an aligned, quantified, one-story application. There is no secret score — there is a gate, a database, and a tired human you can make life easy for.

Sources. Goldman's Oracle tenancy, role data, 4-cap and email rules: pulled directly from hdpc.fa.us2.oraclecloud.com (15 Aug 2026). Workday-to-Oracle switch: Oracle leadership on-record remarks and industry reporting (Cloud Wars; AppsRunTheWorld customer records, 2020). ORC mechanics: Oracle Fusion Cloud Talent Management official documentation — candidate selection processes (phases/states), automated progression conditions, disqualification questions and prescreening, resume parsing/profile import, matching features (Suggested/Similar Candidates) and Recruiting Booster materials; implementation-partner guides (Maverick, Namos, Kovaion) for practitioner detail. GS-specific configuration (question sets, automation rules, whether Booster AI is licensed for campus) is private — everything so marked is labelled BEST GUESS or UNKNOWN, and the live application flow you experience always overrides this document.