How You Are Scored

Who actually watches, what the software does and does not look at, how much each question is worth, what separates a 3 from a 5 on every one of your ten prompts, and whether it is the content or the way you come across.

Doc 31 · Goldman Sachs 2027 EMEA London · Asset Management, Client Solutions Group · Requisition 171037 · Prepared 4 November 2026

CONTENTS

1. **Part I — Who and what is assessing you.** The two possible graders · what HireVue's software actually analyses, in its own words · the myth that will waste your prep · what a human reviewer actually does · BARS, the instrument you are measured against
2. **Part II — Weighting.** The honest position · floor questions versus ceiling questions · my estimated weighting and the reasoning behind it · where the marginal hour of prep goes
3. **Part III — Content or camera, settled with evidence.** What the research actually found · the medium tax · why your voice matters more than your face · the sequential gate model
4. **Part IV — Your ten questions, each with a 1-to-5 scoring anchor.** What it tests, what each score looks like, what makes a 5, the specific trap, and your material
5. **Part V — Delivery, reduced to what moves the rating.** The seven things that count · the things to stop worrying about
6. **Part VI — The self-scoring sheet.** Score your own takes the way a reviewer would

PART I

Who and what is assessing you

1 There are two possible graders, and they want opposite things

Almost everything confusing about HireVue advice online comes from one fact: there are two completely different things that might assess your recording, and they measure different signals. Before anything else, understand both.

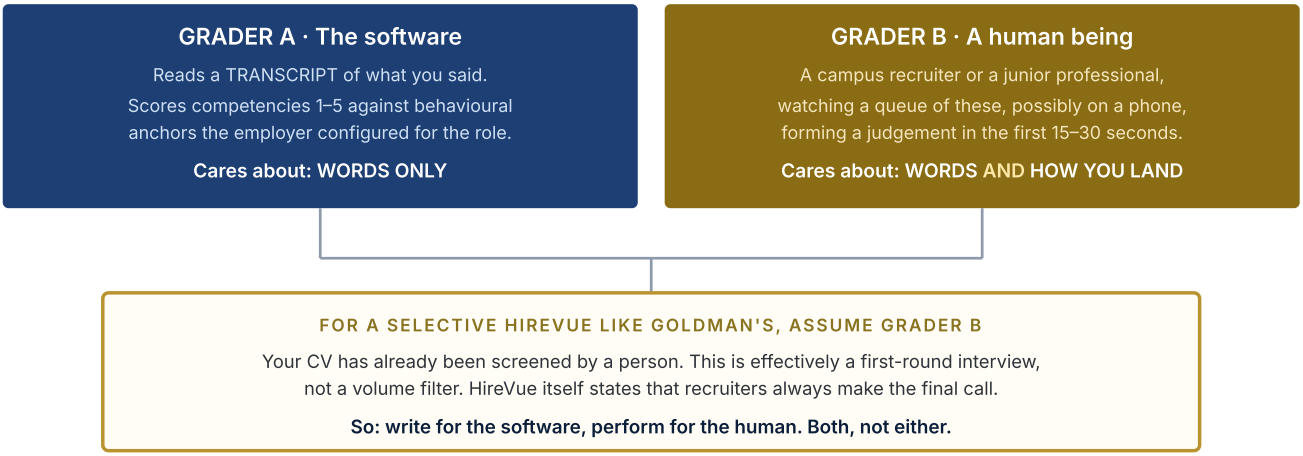


Figure 1 — The two graders. Goldman does not publish which it uses; the reasoning for assuming a human is in section 3.

2 What the software actually analyses — and the myth to drop today

There is a widespread belief, repeated by prep sites and in student group chats, that HireVue scores your facial expressions, eye contact, smiling and tone of voice. **This is out of date and you should stop optimising for it.** HireVue's own science documentation says the opposite, explicitly: PRIMARY — HireVue's published documentation, read 4 November 2026

> "For AI-scored video assessments, Hirevue analyzes the transcript of what a candidate says — not facial expressions, body language, appearance, background, or tone of voice."
>
> "Hirevue does not use facial analysis to evaluate candidates in its assessments."
>
> "Hirevue evaluates candidates using job-relevant skills, behaviors, and competencies defined for a specific role."
>
> "Recruiters always make the final call." · "Hirevue's AI is designed to support human decision-making, not replace it."
>
> HireVue, "Our Science"

Two things follow immediately, and they are the most actionable facts in this document.

1. **If an algorithm touches your recording at all, it is reading a text transcript.** Which means every word has to carry on the page with no delivery attached. A sentence that only works because you said it warmly scores nothing. This is a strong argument for the content-section technique: naming the specific, technical thing in your opening lines reads as well in a transcript as it sounds on camera.
2. **You are not being penalised for not smiling enough, and nothing is reading your face.** The hours students spend worrying about facial expression are being spent against an opponent that no longer exists. Delivery still matters enormously, but for the *human* reason in section 3, which has a different and much more practical set of implications.

How the software scores, if it is used

HireVue's published mechanics:
MEDIUM — HireVue's own documentation and help pages; Goldman's specific configuration is not published

- **Qualification questions** get a binary *Meets* or *Does Not Meet*. These are eligibility facts: graduation date, work authorisation, availability. There is no partial credit and no argument. This is the only part of the process that can hard-reject you regardless of quality.
- **Skill questions** are scored **1 to 5** against a rubric the employer sets, or against one of HireVue's validated competencies, from novice to expert, benchmarked against a norm group.
- The scoring framework is built on **behaviourally anchored rating scales (BARS)**: each score has a written description of what that level of answer actually looks like, so two reviewers rating the same answer land in the same place.

That last point is the single most useful thing to know, and Part IV of this document is built on it. If you know you are being rated against written descriptions of what a 1, a 3 and a 5 look like, then the right preparation is not "be good" — it is "know what a 5 is described as, and produce that."

3 What the human reviewer is actually doing

Goldman's HireVue is selective: your CV cleared a human screen before the invitation was issued. That tells you the firm is not using this as a volume filter and is willing to spend reviewer time on it. The reviewer is most likely a campus recruiter or a junior professional a couple of years in. Assume the following, because it is how this actually works:

- **They are watching a queue.** Possibly dozens. Possibly in a single afternoon. Possibly on a phone.

- **They decide fast.** The working assumption across every practitioner account is that the advance/no-advance instinct forms inside the first fifteen seconds, and attention degrades sharply after thirty. You cannot fully win in fifteen seconds. You can absolutely lose there.
- **They form one overall impression, then rationalise it.** This is the part nobody tells candidates. In practice a reviewer forms a single gestalt judgement — *would I want this person sitting next to me* — and the per-question scores tend to line up behind it. That means consistency across answers matters more than peaking on one. A brilliant market-news answer attached to four limp ones does not average out; it reads as a rehearsed party piece.
- **The bar is lower than you fear.** Most candidates visibly improvise. A structured, rehearsed answer stands out not because it is extraordinary but because the comparison set is weak.

> **THE REVIEWER'S REAL QUESTION**
>
> They are not asking "is this answer correct." They are asking **"can I put this person in front of a client, or next to one of our investment people, without worrying?"** Every question in your set is a different angle on that one question. Once you see that, the behavioural questions stop looking like filler: they are the "without worrying" half.

PART II

Weighting

4 The honest position on weighting

> **READ THIS BEFORE THE TABLE**
>
> Goldman does not publish a weighting, and neither does anybody else. Any prep site giving you exact percentages is making them up. **So is this document — the difference is that I am telling you.** What follows is my reasoned model, built from: which questions are reported for GS Asset Management, what each one can actually discriminate between candidates on, the fact that your invitation explicitly promises division-specific questions, and how selective HireVues are generally used. Treat it as a *prep allocation tool*, which is what it is good for, and not as a disclosed rubric, which it is not.

The distinction that actually matters: floor questions and ceiling questions

Percentages are the wrong mental model anyway. The useful one is **asymmetry**. Your ten prompts split cleanly into two groups that behave completely differently.

	FLOOR questions (behavioural)	CEILING questions (soft-technical)
Which	Teamwork · strength and weakness · something not on your CV · what motivates you	Walk me through your CV · why Goldman · why Asset Management / CSG · recent market news · a challenge · why should we hire you
What they do	Confirm you are a normal, pleasant, self-aware human who can be put in a room. Pass/fail in nature.	Separate you from other candidates who are also normal, pleasant and self-aware.
Best realistic outcome	"Fine. Next." You cannot win here.	"I want to meet this person."
Worst realistic outcome	**Rejection.** A weird, boastful, evasive or hollow answer here kills an otherwise strong recording.	"Generic, like everyone else." Damaging but survivable if the rest is strong.
Prep strategy	Make them *safe* and *human*. Thirty minutes each, maximum. Do not over-engineer.	Make them *un-copyable*. This is where the hours go.

The practical consequence: **you cannot win the offer on the behavioural questions, but you can lose it there.** The most common sophisticated mistake is a candidate who pours everything into "why Goldman" and then gives a stilted, over-finance-ised answer to "something not on your CV" that makes the reviewer think the whole thing was a performance.

5 My estimated weighting, with the reasoning

MY MODEL — NOT A DISCLOSED RUBRIC Percentages are share of the advance/no-advance decision, assuming four to six questions are drawn from this bank. "Upside" is how much a brilliant answer gains you; "downside" is how much a poor one costs you. Note that they are not symmetric, and that asymmetry is the real insight.

Question	Weight	Upside	Downside	Why I weight it there
Walk me through your CV	20%	High	High	Near-certain opener, so it sets the gestalt impression every later answer is read against. It is also the purest test of whether you can structure and prioritise. Both upside and downside are large, which is rare.
Why Asset Management / this division	20%	High	High	Your invitation explicitly promises questions from each business you applied to, so this is close to guaranteed. A divisional reviewer cares more about whether you understand *this seat* than about whether you admire the firm. Most candidates are much weaker here than on "why Goldman".
Recent market news	15%	Very high	Medium	The biggest spread between a good and a bad answer of anything in the set, because the average standard is so low. The lower downside reflects that a mediocre one reads as "student", not as "problem".
Why Goldman Sachs	15%	Medium	High	High downside, capped upside. Almost everyone prepares it, so a good answer is expected rather than impressive; a generic one is actively damning because it signals you did not do the work.
A challenge you faced	10%	Medium	Medium	Sits between the two groups. Treated as behavioural by most candidates, which is exactly why putting an analytical spine in it is cheap differentiation.
Why should we hire you	5%	Medium	Medium	Less commonly reported. If it appears it is a synthesis question: a chance to restate your three strongest pieces of evidence without adjectives.
What motivates you	5%	Low	Medium	No target answer exists, so honesty plus one concrete anchor is the whole game.
Teamwork / leading others	4%	Low	High	Pure floor. The downside is high because a vague or self-aggrandising answer here is a real red flag for a client-facing seat.
Strength and weakness	3%	Low	High	Pure floor, high downside. A fake weakness is one of the few things that can single-handedly sink a recording, because it reads as either unselfaware or dishonest.
Something not on your CV	3%	Low	High	Pure floor. Only testing whether you are an interesting human. Smuggling finance in here is the classic own goal.
Total	**100%**			**Ceiling questions ≈ 70% · Floor questions ≈ 30%, but with the highest concentration of fatal downside.**

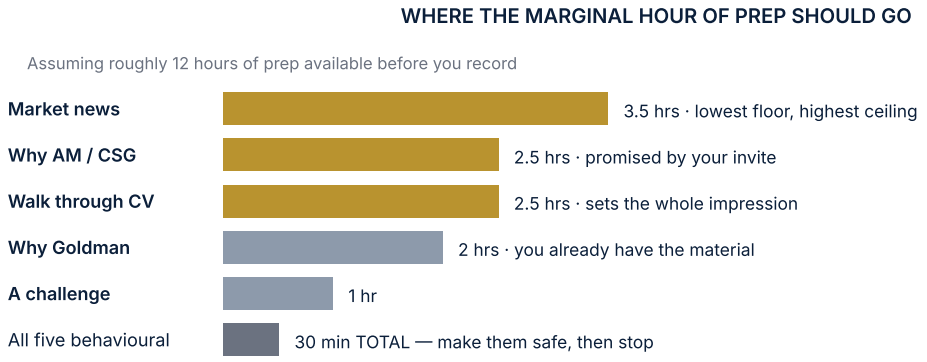


Figure 2 — Prep allocation implied by the model. MY MODEL

THE ONE PLACE THE MODEL BREAKS

Thirty minutes total on the behavioural five assumes you have honest, specific answers already available in your head. If you do not — if "a time you worked in a team" produces a blank — then that question is temporarily your highest-value prep, because an unprepared floor question is the single most likely way this recording goes wrong. Check you can answer all five out loud right now. If any produces a blank, fix it before touching the market-news answer.

PART III

Content or camera, settled with evidence

6 What the research actually found

You asked whether it is what you say or the way you come across. This is answerable with evidence rather than opinion, and the evidence points somewhere more useful than "both matter."

Finding	What it means for you
Recorded and technology-mediated interviews produce systematically *lower* ratings than face-to-face ones. A reported meta-analytic effect size of about d = −0.41 on interviewer ratings. MEDIUM — from a research abstract, full methods not verified	There is a **medium-sized penalty built into the format itself**. The same person, giving the same answer, scores worse on camera than in a room. This is not about you. It is the single most important delivery insight available, and the implication is in section 7.
In thin-slice studies of interviews, *audio* cues predicted hireability ratings more accurately than *visual* cues. Visual cues were comparatively poor predictors. MEDIUM — Nguyen & Gatica-Perez, ICMI 2015	**Your voice matters more than your face.** Pace, energy, pitch variation and fluency are doing more work than eye contact or expression. If you have limited rehearsal time, rehearse out loud rather than in front of a mirror.
In classic videotaped-interview research, nonverbal cues accounted for roughly 58–62% of the variance in judges' ratings. LOW — older study, single source, dated methodology	Treat the exact number as unreliable, but the direction is consistent with everything else: **how you deliver carries a large share of a human reviewer's rating**, whether or not they would say so if asked.
Judges infer *social skill* from nonverbal behaviour far more accurately than they infer *motivation*. LOW — same older study	Delivery communicates "would be good with clients." It does *not* communicate "actually wants this job." Motivation has to be carried by **content**: specific, verifiable, un-copyable evidence. You cannot sound your way into being believed about why you want it.
HireVue's AI, where used, reads the transcript only. PRIMARY	Any algorithmic layer is pure content. Delivery is a human-channel phenomenon exclusively.

7 The answer: they are sequential gates, not competing weights

The question "is it content or delivery" assumes they trade off against each other. They do not. They operate at different moments and they fail in different ways.

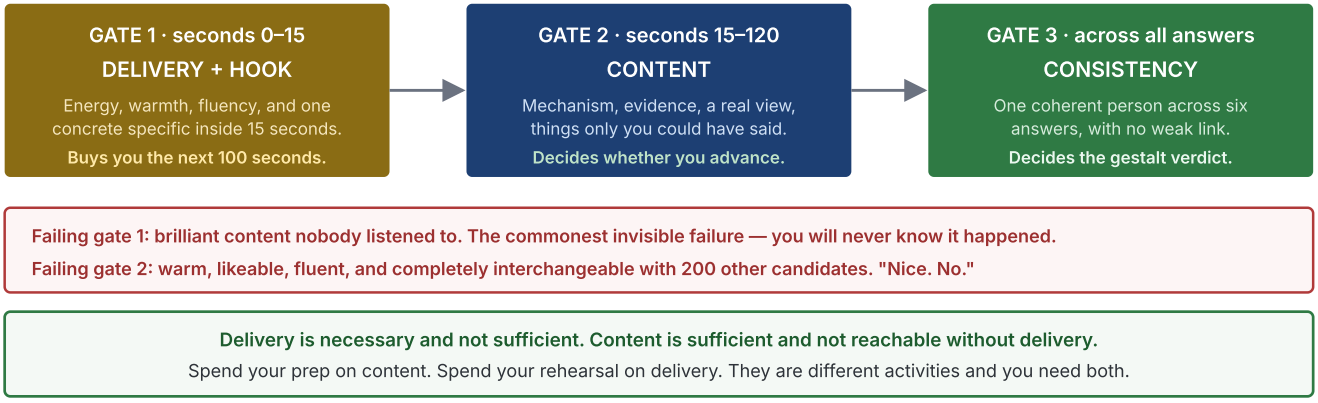


Figure 3 — The sequential gate model. MY SYNTHESIS of the evidence in section 6

The medium tax, and what to do about it

This deserves its own paragraph because it is the most actionable research finding here. If recorded interviews systematically produce lower ratings than live ones, the mechanism is almost certainly that **the format strips out the things that make a person seem engaging**: there is no listener nodding back at you, no interruption, no laughter, no feedback loop. Your natural conversational energy is calibrated against another human reacting. Record into a webcam and that calibration produces something flat.

The correction is to deliberately overshoot. Speak with noticeably more energy, more pitch variation and slightly more animation than feels natural to you. It will feel like about 15 to 20% too much. On the recording it will land as normal. Every candidate who has watched their own takes back has had this exact realisation; getting there before you record is what the rehearsal is for.

PART IV

Your ten questions, scored

8 What a 1, a 3 and a 5 look like on each prompt

Since HireVue scoring is built on behaviourally anchored rating scales, the most useful possible preparation is to see the anchors. **Goldman's actual anchors are not published, so these are mine** MY CONSTRUCTION — written from what each question can discriminate on and what reviewers consistently report rewarding. They are deliberately written as descriptions of answers, not as adjectives, so you can score your own takes against them in Part VI.

CEILING 20%

Walk me through your CV / tell us about yourself

What it is really testing: can you prioritise? A CV has twenty things on it and you have two minutes. The choice of which three to keep *is* the answer. Secondarily, it tests whether your path has a logic or is a list of accidents.

1	Chronological recital from school onwards. Equal time to everything. Runs out of time before reaching the relevant experience. No through-line.
3	Covers the right experiences in sensible order with reasonable fluency. Says what they did but not what it produced or what they learned. Could be anyone with a similar CV.
5	Opens by naming the three things they will cover, with a specific artefact or figure in the first fifteen seconds. Each one connects to the next with a reason, so the path reads as a decision rather than a sequence. Ends pointing directly at this role. The listener could repeat the three threads back afterwards.

The trap: chronology. Chronology feels safe and guarantees a 2. The strongest version is thematic, not time-ordered.

Your material: three threads — Pictet Global Client Group (the client-facing side of an institutional asset manager, where you built the DDQ portal across roughly 400 lines and 20 teams); the fifty segregated market openings (the operational work that turns a won mandate into a live portfolio); and the Goldman spring week (why this firm specifically, not in the abstract). The through-line that connects all three honestly: each one was about knowing the whole of something well enough to connect the right part of it to the person who needed it.

CEILING 20%

Why Asset Management / why this division?

What it is really testing: do you know what the job actually is, as opposed to what the division is called? This is the question where candidates who researched the *firm* but not the *seat* are exposed, and it is close to guaranteed to appear because your invitation promises division questions.

1	Describes Asset Management in words that would apply to any asset manager, or worse, describes investment banking. Shows no awareness that Client Solutions Group is client-facing rather than investing.
3	Correctly describes what asset management does and gives a plausible personal reason for the interest. Mentions long-term client relationships. Nothing a reviewer has not heard four times that morning.
5	Names what the seat actually does — sitting between an institutional client and the firm's investment capability and translating in both directions — and evidences having done a version of it. Knows which revenue line the function feeds and why the firm is investing there. Distinguishes this division from the investing seats deliberately rather than by accident.

The trap: describing the *investing* job. CSG does not manage portfolios. A candidate who talks about wanting to pick investments has applied to the wrong requisition and the reviewer will know within one sentence.

Your material: you have already done the function. The DDQ portal is the artefact an institutional client needs in order to defend a decision to its own board; the market openings are the onboarding that follows a won mandate. And the firm-level fact that makes it an argument rather than a preference: management and other fees are the line this function feeds, they hit a record and grew 22% year on year, and the firm has been deliberately shifting toward durable fee revenue since the consumer episode. You would be joining the part of the firm the firm is betting on.

CEILING 15% A recent piece of market news / something in the world that interests you

What it is really testing: can you think in cause and effect, and do you actually care? Not what you know. See doc 29 for the full treatment.

1	Lists several topics with one sentence each, or summarises one article accurately and stops. Every thought in the answer belonged to the journalist. Ends on "it will be interesting to see what happens."
3	One story, explained correctly, with a stated opinion. The opinion is a relayed consensus view or a directional forecast with no mechanism behind it. Competent and forgettable.
5	One story. States the mechanism as a chain the listener can follow. Names a specific type of investor and a specific decision that is now harder for them. Gives a first-person claim with a reason, volunteers the counter-argument, and names what would prove them wrong. Every figure attributed.

The trap: breadth. Rates, gold, inflation and a forecast is a 1, not a 4, however much ground it covers.

Your material: the AI data-centre debt chain, from chip order to pension fund in four steps, ending on the observation that the exposure has become hard to *see* rather than hard to judge. Backup: the OCIO and defined-benefit surplus story, where the Verizon and Lockheed mandates are the symptom and the structural shift is the disease.

CEILING 15% Why Goldman Sachs?

What it is really testing: did you do work that the other candidates did not? High downside, capped upside — everyone prepares this, so a good answer is the expectation.

1	Prestige, reputation, "the best people", "leading global institution", "I've always admired". Swap in any bank's name and the answer survives unchanged.
3	Names something real about the firm — the scale, the breadth, a recent deal — and connects it loosely to a personal interest. Accurate, researched, and still reachable from the careers site in twenty minutes.
5	Makes a structural argument about why this firm's model produces something a competitor's cannot, supports it with a figure, and attaches a first-hand experience that is theirs alone. Would collapse if you substituted another bank's name, which is the test.

The trap: adjectives. "Collaborative", "entrepreneurial", "client-focused" are what the firm says about itself, so repeating them proves only that you read it back.

Your material: the structural argument is that the asset management arm sits inside a firm that originates private deals, finances institutional investors and advises corporates, so the capability you would be representing is partly manufactured in-house rather than bought — evidenced by alternatives assets up 26% year on year to a record $491bn. The first-hand part is yours alone: you heard the integration argument made in person, by Ahmed on consolidating competing internal teams onto one platform, and by Jack describing the client-facing seat as the conductor of the orchestra. You were in the room. Almost nobody else was.

MIXED 10% A challenge you faced — how did you approach it, what did you learn?

What it is really testing: nominally resilience. Actually, whether you have a method. Most candidates treat it as behavioural, which is exactly why an analytical spine is cheap differentiation here.

1	A challenge that was not really difficult, or a difficulty with no action attached. The "lesson" is a platitude: I learned to manage my time, I learned the importance of communication.
3	A genuine challenge with a clear situation, action and outcome. Properly structured. The lesson is reasonable but generic, and the approach could have been anyone's.
5	A real constraint, named precisely. A diagnosis of *why* it was hard before any action is described. A method that was chosen for a reason, with the trade-off acknowledged. A specific outcome. A lesson that is about judgement rather than effort, and that the person has visibly applied since.

The trap: choosing a challenge where the solution was "I worked harder". That answers an endurance question nobody asked.

Your material: the DDQ portal. The constraint was real and diagnosable — answers were owned by twenty different teams, none of whom reported to you, and the client's deadline was fixed. The interesting part is not that you finished it; it is that you had no authority and therefore had to solve it with structure instead. That is exactly the shape of the CSG job, which is why this is the right story for this application.

MIXED 5% Why should we hire you?

What it is really testing: can you select evidence? It is a synthesis question and the temptation to list qualities is almost irresistible.

1	A list of adjectives: hard-working, passionate, a fast learner, a team player. No evidence attached to any of them.
3	Two or three claims, each with an example. Reasonable. Reads as a summary of the CV already discussed.
5	Names one trait, not four, and evidences it from two or three genuinely independent contexts so it reads as a pattern rather than an anecdote. Connects the trait explicitly to what the seat requires. Short and unhurried.

Your material: one trait — knowing the whole of something well enough to connect the right part of it to the person in front of you — shown three times in increasingly institutional settings: the shop, where you learned it without a name for it; the DDQ portal, across twenty internal teams; the market openings, across operational counterparties. Same trait, three independent contexts. Very few applicants can do that.

FLOOR 5% What motivates you / what drives you?

What it is really testing: whether you are self-aware and whether the motivation is consistent with the job. There is no target answer, which is precisely why people invent one.

1	An answer chosen because it sounds impressive: "I'm driven by challenge", "by excellence", "by being the best". Nothing concrete attached.
3	An honest motivation, stated plainly, with a general gesture at where it came from.
5	An honest motivation anchored immediately to one specific moment where it was visible, and a brief note of why it fits this kind of work. Sounds like a person describing themselves, not a candidate answering a question.

The trap: performing an ambitious-sounding motivation. Reviewers hear the difference, and an honest ordinary motivation beats an impressive invented one every time.

Your material: you have said yourself that you are motivated by helping others get what they actually need, and that it goes back to the shop — the satisfaction was never the till, it was knowing the shelf well enough to link one product to the next until the customer had what they came for. That is honest, it is specific, and it happens to describe the job.

FLOOR 4% Tell me about a time you worked in a team / led others

What it is really testing: whether you can share credit and whether you understand what other people were doing. For a client-facing seat the downside here is higher than the weight suggests.

1	"I" throughout. The team is scenery. Or the reverse: so much "we" that no individual contribution is visible. Any hint of blaming teammates is close to disqualifying.
3	A real team situation, a clear personal contribution, a decent outcome. Appropriate balance of I and we.
5	Shows awareness of what the other people needed, not just what the task needed. Describes adapting to a specific person or a specific friction. Credits someone else by name or role without performing humility. Leadership, if claimed, is shown through a decision rather than a title.

The trap: a group coursework story told as though you carried everyone. It is the most common answer and the least attractive.

FLOOR 3% A strength and a weakness

What it is really testing: honesty and self-awareness, in that order. One of the few single answers that can sink a whole recording.

1	A humble-brag: "I'm a perfectionist", "I care too much", "I work too hard". Or a weakness so trivial it is obviously evasive. Reads as either unselfaware or dishonest, and both are fatal for a client-facing role.
3	A real weakness, admitted plainly, with something being done about it. Safe and adequate.
5	A real weakness, named without drama, with a concrete instance where it cost something, and a specific mechanism now used to manage it rather than a vague intention to improve. Delivered without apology or over-explanation, which is itself the evidence of self-awareness.

The trap: the perfectionist answer. Every reviewer has heard it hundreds of times and it reads as a refusal to answer the question.

FLOOR 3% Something not on your CV / your interests

What it is really testing: whether you are an interesting human being that somebody would want to sit next to. That is the entire test. Nothing else.

1	Smuggles finance in: a personal portfolio, a trading app, a markets newsletter. Or appends a justification — "and I think this shows my discipline" — which signals that the candidate does not trust the answer to stand alone.
3	A genuine non-professional interest, described pleasantly. Reading, football, travel. Fine. Forgotten.
5	Something specific enough to be memorable and to have a detail in it that only a real participant would know. Told with some warmth and ideally one light moment. No justification appended, because none is needed.

The trap: the appended justification. "This shows my grit" turns an answer about a human being into another answer about a candidate.

Note on the light moment: this question is the single best place in the recording to let the reviewer see you enjoy something. A small genuine smile here does more for the gestalt verdict than any amount of polish elsewhere, because it is the only question where it is unambiguously appropriate.

PART V

Delivery, reduced to what moves the rating

9 Seven things that count, and four that do not

Given that any algorithmic layer reads only the transcript, everything below is about the human reviewer. Ordered by how much I think it actually moves a rating. MY ORDERING, from the evidence in Part III

#	What	Why it matters and what to do
1	**Energy, overshot**	The format drains perceived engagement. Speak with roughly 15–20% more energy and pitch variation than feels natural. It will feel excessive and will land as normal. This single adjustment probably does more than everything below it combined.
2	**A specific in the first 15 seconds**	Not an adjective. A named thing, a figure, an artefact. This is what buys the remaining hundred seconds, and it works on both graders simultaneously.
3	**Fluency without a script**	Audio cues predict ratings better than visual ones, and the strongest audio signal is unhesitant speech that is clearly not being read. Memorise four beats and three figures, generate the words live. Film fifteen takes deliberately changing the words each time.
4	**Pace and the deliberate pause**	Nerves accelerate speech. A short pause before the most important sentence does two things: it signals that what follows is the point, and it reads as composure. Most candidates have no pauses at all.
5	**Warmth at the open and close**	Thank them for watching. They did not want to be there. It costs two seconds and it is the cheapest differentiation available.
6	**Camera at eye level, lit from the front**	A laptop on a desk films up your nose and makes anyone look worse. Raise it on books. Face a window or a lamp; do not sit with a window behind you. Plain background. This is about removing distractions, not gaining points.
7	**Looking at the lens, not the screen**	Eye contact is simulated by looking at the small camera hole, not at your own image. If you use notes, keep them to two or three words a line, directly under the lens. A full script is obvious because the eyes track left to right.

AND FOUR THINGS TO STOP SPENDING ENERGY ON

- **Facial expression analysis.** HireVue states it does not do this. Stop managing your face and manage your voice instead.
- **Keyword stuffing for an algorithm.** If a transcript is scored, it is scored against behavioural anchors for competencies, not against a keyword list. Jargon with no mechanism behind it reads as thin to both graders.
- **The perfect take.** A small stumble you recover from naturally is invisible. Restarting six times until it is flawless produces something that sounds recited, which is far more damaging.
- **Clothing beyond a suit and tie.** Wear one. There is no upside to anything else and no further marks available.

> **ON THE RETAKE QUESTION**
>
> Reported retake policies for Goldman genuinely conflict between sources and I could not resolve it. **Prepare as though you get exactly one take on every question.** If the platform does offer a retake when you get there, that is a bonus — and note that the decision screen between attempts typically has no timer, which is the real value of a retake, not the second attempt itself. But do not build your plan on it.

PART VI

Score your own takes

10 The self-scoring sheet

Print this. Record a take, watch it back, and score it honestly before re-recording. Scoring your own answers against written anchors is the closest you can get to seeing what the reviewer sees, and it is far more useful than simply filming again and hoping.

Criterion	1	2	3	4	5
First 15 seconds contain a named specific, not an adjective	☐	☐	☐	☐	☐
Structure is audible — a listener could name the parts afterwards	☐	☐	☐	☐	☐
Originality — would a well-prepared competitor say these same sentences?	☐	☐	☐	☐	☐
Evidence — claims are attached to things that actually happened	☐	☐	☐	☐	☐
Reasoning — there is a chain, not a list	☐	☐	☐	☐	☐
Energy — watching it back, does it sound engaged or flat?	☐	☐	☐	☐	☐
Fluency — spoken, not recited; no visible reading	☐	☐	☐	☐	☐
Length — lands between 90 and 120 seconds with a buffer	☐	☐	☐	☐	☐
Warmth — thanked them; sounds like a person, not a submission	☐	☐	☐	☐	☐
Accuracy — every figure verified and attributable	☐	☐	☐	☐	☐

How to read your score. Under 30, re-record. 30 to 40, re-record once more targeting your two lowest rows. Over 40, submit it and move to the next question. And one override rule: **if "Originality" scores 2 or below, the total does not matter.** An answer that any well-prepared competitor could also give is the defining failure mode of this format, and no amount of polish on the other nine rows compensates for it.

> **THE SINGLE SENTENCE TO TAKE AWAY**
>
> Content decides whether you deserve to advance; delivery decides whether anyone finds out; and consistency across all six answers decides what they actually remember. Prepare the content, rehearse the delivery, and do not let any one answer be the weak link — including the ones that look like filler.

Sourcing and confidence

Primary. HireVue's statements about what its assessments analyse, the absence of facial analysis, the role of human reviewers and the use of behaviourally anchored rating scales come from HireVue's own published science documentation, read on 4 November 2026 and quoted verbatim.

Medium. The 1-to-5 competency scale, the Meets / Does Not Meet treatment of qualification questions and the norm-group benchmarking come from HireVue's help documentation and secondary summaries; Goldman's specific configuration is not published anywhere. The research findings on technology-mediated interview ratings ($d \approx -0.41$) and on audio cues out-predicting visual cues come from research abstracts and a conference paper whose full methods I did not verify. The reported question set is from student and forum reports, not from Goldman.

Low. The older videotaped-interview study's claim that nonverbal cues explain 58–62% of rating variance rests on a single dated source; I have used it only for direction, not magnitude.

Mine, and labelled as such throughout. The weighting table, the floor-versus-ceiling framing, the prep allocation, the sequential gate model, the ordering of delivery factors and every one of the 1-to-5 behavioural anchors in Part IV are my construction. Goldman's real rubric is not public. These are built to be a useful preparation instrument, and they will be directionally right and specifically wrong in ways I cannot see. Use them to allocate effort and to score your own takes, not as a claim about what is on a reviewer's screen.

Unresolved. Whether Goldman permits retakes, and how many questions you will actually get. Prepare for one take and for the full set.