The Four AI Answers

Your Pictet story rebuilt around the two things you have not noticed in it, plus the accuracy question, the forward-looking question, and your day-to-day use — all four, written out.

Doc 50 · Supersedes the drafts in doc 47 · Asset Management, Client Solutions Group · Prepared 5 November 2026

PART I

The Pictet story

1 Three things you are underselling

The detail you have given me now is much better than the earlier version. But you are describing it as "I automated a process", and three things in it are considerably more impressive than that.

1 · YOU DIAGNOSED IT AS A ROUTING PROBLEM, NOT A WRITING PROBLEM

Anyone looking at "four to five hundred due diligence questions" assumes the bottleneck is **answering** them. You noticed it was **getting each question to the right person.**

And you quantified the real cost correctly: a misrouted question was not one wasted email, it was roughly **five** — send, wrong person, bounce back, reassign, forward. *The answer was never the expensive part. The ping-pong was.*

That is the diagnosis beat, and it is the single highest-scoring thing in the story. Most candidates describe what they built. Almost none describe how they worked out what to build.

2 · YOU DESIGNED YOURSELF OUT OF THE LOOP

Look at what actually changed:

Before: question arrives → the team routes it → wrong expert → comes back to the team → the team reassigns → right expert.
After: question arrives → auto-assigned → if wrong, *that person reassigns directly.*

You removed the middleman. **And the middleman was partly you.**

That is unusual and worth saying out loud. Most people build tools that make themselves more central. You built one that made your own coordinating role unnecessary — which is what somebody does when they are solving the team's problem rather than demonstrating their own value.

3 · YOU CHOSE A USE CASE WHERE BEING WRONG WAS CHEAP — AND THIS IS THE BEST BIT

You did not point AI at writing the answers. **You pointed it at deciding who should write them.**

Think about the asymmetry. If the routing is wrong, a human notices in thirty seconds and reassigns — *the cost of error is near zero.* If AI had drafted a DDQ answer and it was plausible but wrong, that is **a representation the firm is making to a client**, used to justify an allocation to a board, and it might not surface for years.

You instinctively put AI on the high-frequency, low-consequence task and kept humans on the low-frequency, high-consequence one. That is exactly the judgement a regulated firm wants to see, and you arrived at it without being told.

Say this. It is the thing that turns a tool story into a judgement story.

2 The arithmetic you can do live

You do not have a clean before-and-after number, and inventing one is not an option. **But you can do the sum out loud, which is actually better** — it shows the reasoning rather than asserting a result.

THE MOVE

"Five hundred questions a cycle. *If even one in five went to the wrong person first* — and it was more than that — that's a hundred misroutes, at about five emails each. **So five hundred emails of pure ping-pong, before anybody had answered anything.**"

Note the construction: you are not claiming a measured figure. You are showing the order of magnitude from numbers you genuinely know, and flagging your own assumption. **That reads as more honest than a precise number, not less.**

Still worth ten minutes tonight: check whether you can find the actual misroute rate, or how long a DDQ cycle took before and after. A real figure beats a worked estimate. But the estimate is a perfectly good fallback and you should not stress about it.

3 Answer 1: a time you used AI

About 115 seconds.

CONTENT SECTION — ~15 SECONDS

Thanks for taking the time to watch this. I'll talk about a problem I found on my placement at Pictet, in the client-facing team — how we handled due diligence questionnaires, which is the long list of questions institutional clients send about the mandates they have with you. And the thing I'd most want to get across is how I worked out what the actual problem was, because it turned out not to be the one everyone assumed.

THE PAIN POINT — ~25 SECONDS

So we'd get four to five hundred questions in a cycle, and they were being tracked on an Excel tracker that hadn't properly been updated in years, with everything else happening over email. And people's inboxes were getting three or four hundred emails a day.

The specific thing that struck me was what happened when a question went to the wrong expert. It'd go out, come back to us, we'd reassign it, send it on again. So one misrouted question was about five emails before anyone had even started answering it.

THE DIAGNOSIS — ~25 SECONDS

And that's when I realised the bottleneck wasn't writing the answers — it was routing them. Everyone had assumed the problem was volume, because four hundred questions sounds like a lot of writing. But the experts could answer their own questions quite quickly. The expensive bit was the ping-pong getting each question to the right desk.

If even one in five went to the wrong person first — and it was more than that — that's a hundred misroutes at five emails each, so five hundred emails of pure admin before anybody had answered anything.

WHAT I BUILT, AND THE DESIGN CHOICE — ~30 SECONDS

So I used Claude Code to build a dashboard, in a sandboxed environment so we weren't putting anything at risk while we tested whether it worked. It automatically assigned each incoming question to the team most likely to own it — and crucially, if it got that wrong, the person who received it could reassign it directly rather than sending it back to us. So it cut the middleman out, and the middleman was partly me.

And the thing I'd point to as the actual judgement call is what I pointed it at. I didn't have it write the answers. A DDQ response is effectively a representation the firm is making, and a client uses it to justify an allocation to their board — so if that's plausible but wrong, you might not find out for years. Whereas if the routing's wrong, somebody notices in thirty seconds and fixes it. So I put it on the high-frequency, low-consequence bit and kept people on the other one.

RESULT AND LIMIT — ~20 SECONDS

The team were getting through questions in a much more streamlined way, and we escalated it upward to try to get it approved properly at firm level. [Say exactly how far it got.]

And the honest limit is that it didn't make anybody's answers better. It just stopped the time being spent on getting the question to the person who already knew the answer. Which I think is usually where the real waste is.

PART II

4 The risks, ranked by how much they actually matter

You said you were not sure on this one. Here they are, ordered — and **the first and the last are the ones worth saying.**

#	Risk	Why it matters here
1	**Plausible but wrong**	**The big one.** The dangerous output is not the obviously wrong one — it is the fluent, confident, correctly formatted one that is subtly incorrect. Obvious errors get caught. Plausible ones get used.
2	**Confident about specifics it cannot know**	Figures, citations, quotes, dates. The exact things that look most authoritative are the ones most worth checking.
3	**Stale**	Training data has a cutoff. In markets, a six-month-old fact is often just wrong.
4	**No provenance**	You cannot always trace where an assertion came from — which matters enormously when the output becomes something the firm is asserting to a client.
5	**Confidentiality**	Putting client material into a tool that is not approved or not contained. The reason Goldman runs approved models inside its own environment.
6	**Automation bias**	**The subtle one, and the best thing you can say.** When something is usually right, people under time pressure stop checking. The risk is not the tool failing — it is the humans quietly ceasing to verify because it has been fine for six months.
7	**Homogenisation**	If every firm uses similar tools, outputs converge — awkward in a business where differentiation is the point.

5 Two principles that answer "how do you make it accurate"

Principle one: **it is reliable in inverse proportion to how much you need it.** If you know enough to check the answer, it is extremely useful. If you do not, you cannot tell whether it is right — which is exactly when you are most tempted to trust it.

That is why the right use is in areas where you *already* have some expertise, and why "I use it to learn things I know nothing about" needs the verification step bolted on, which is what your reading workflow does.

Principle two: **match the check to the consequence.** Routing a question wrong costs thirty seconds. A wrong client-facing answer costs years. So the verification should be proportionate — light where error is cheap, full human sign-off where it is not.

And you did not read that in an article. **You applied it at Pictet without naming it**, which is why this answer and your first answer reinforce each other.

Answer 2: how do you make sure it is accurate

ABOUT 80 SECONDS — THIS ONE SHOULD BE SHORTER

Thanks for watching. The thing I'd say first is that the risk isn't an output that's obviously wrong — those get caught. It's the one that's fluent and well-formatted and subtly incorrect, because that's the one people use.

So the rule I've landed on is to match the checking to the consequence. When I built the DDQ dashboard at Pictet, I deliberately pointed it at routing questions rather than writing answers. If the routing's wrong, somebody notices in thirty seconds. If a client-facing answer is wrong, that's a representation the firm is making, and you might not find out for years. So the high-consequence bit kept a person on it.

And for my own research I've got a simple rule, which is that I don't quote anything I haven't read in the original. Summaries drop the caveats, and the caveats are usually where the interesting part is.

I think the risk people underrate is the slow one — when a tool is usually right, you gradually stop checking, because it's been fine for six months. So the discipline isn't really about the technology, it's about not letting the verification quietly stop.

PART III

6 First — why your prospecting idea is the wrong answer

GENTLY, BUT CLEARLY

You suggested AI will help find contacts, scrape LinkedIn, identify pension funds. **I would drop this entirely**, for three reasons:

1. It solves a problem that does not exist. Everyone already knows who the UK pension schemes are. There is a published universe, consultants track it, and the data is purchasable. *Prospect identification has never been the bottleneck.*

2. It contradicts the trigger model. From doc 48 — you do not win clients by finding them, you win searches when something triggers one. Knowing a scheme exists is worth nothing; being credible when they look is worth everything.

3. Scraping carries real compliance issues. Mentioning it in a regulated-firm interview is a small unforced risk you do not need to take.

7 What actually changes — a three-link chain

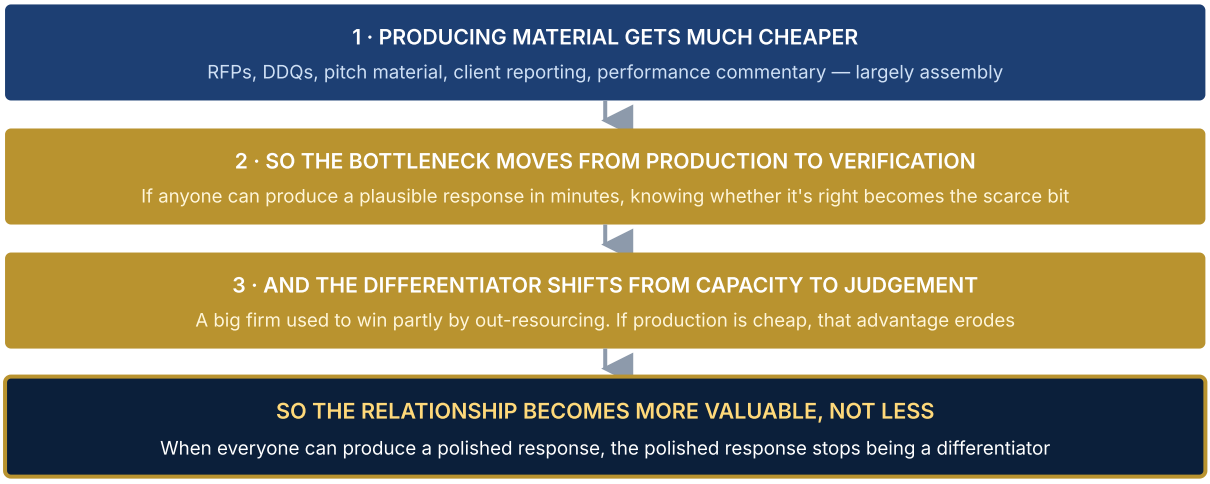


Figure 1 — The chain. My reading, not a published view — hold it as a position you can defend.

WHY THE CONCLUSION IS THE INTERESTING BIT

The obvious prediction is that AI makes client-facing roles *less* necessary. **I think the opposite follows, and so should you.**

What cannot be commoditised: knowing which board member is sceptical and why. Turning up after a bad quarter with an honest explanation. Knowing what a client needs before they have articulated it. **AI commoditises the output and therefore raises the premium on everything it cannot produce.**

Three more you can draw on

- **The buyer gets AI too.** A trustee board can now summarise and compare manager responses far faster than before — which means *weak answers get caught more easily.* The standard rises on both sides. Almost nobody thinks of this.
- **Client expectations rise.** Faster turnaround, more customisation, more frequent reporting. The floor moves up for everyone, so speed stops being a differentiator and becomes table stakes.
- **What it cannot do is be accountable.** A client is not buying the document — they are buying a named institution standing behind it. AI can draft the answer; it cannot be responsible for it.

Answer 3: how will AI change asset management and this role

ABOUT 95 SECONDS

Thanks for watching. I'd start from what's actually expensive in a client-facing asset management job, which is producing material — RFPs, due diligence questionnaires, client reporting, performance commentary. And a lot of that is assembly from things the firm has already said and approved, which is exactly the kind of work this is good at. I saw a version of that myself at Pictet.

But I think the more interesting consequence is the second-order one. If producing a plausible response gets cheap, then the bottleneck moves from production to verification. Knowing whether something is right becomes scarcer than being able to write it. And the differentiator shifts with it — a large firm used to be able to win partly by out-resourcing everyone, having more people to respond to more searches. If production is cheap, that advantage erodes.

Which leads me somewhere slightly counterintuitive. You'd assume AI makes client-facing roles less necessary. I think it does the opposite. When everyone can produce a polished response, the polished response stops being what distinguishes you. What can't be commoditised is knowing which trustee on the board is sceptical and why, or turning up after a bad quarter with an honest explanation. So it raises the premium on exactly the part that was always the actual job.

And the thing I'd add is that the clients are using it too — a trustee board can compare manager responses much faster now, which means a weak answer gets caught more easily than it used to. So the standard goes up on both sides at once.

PART IV

Day to day

8 Your two personal uses, and what connects them

You gave me two: the Gmail colour-coding, and the research-to-audio workflow. **There is a single idea running through both and through your Pictet project, and you have not named it.**

What you did	What AI did	What you kept
Pictet DDQ	Routed the question to a likely owner	Answering it. And overriding the routing.
Gmail triage	Classified mail — needs action, CC'd, has a deadline	Deciding what to do and in what order
Research workflow	Gathered and converted to audio for a baseline	Choosing what to go deep on, and verifying it

In all three, **AI does the sorting and you do the deciding.** That is not three tools — it is one consistent philosophy, and saying so is far stronger than listing three examples.

AND THE SPECIFIC THING THAT IS GOOD ABOUT YOUR LEARNING WORKFLOW

You said it almost in passing: you listen at the gym to get a baseline, **and then go back and read the sections you did not understand.**

That is not consuming content. **You are using the audio pass to find out what you do not know, then reading to fix it.** The first pass is diagnostic; the second is the actual learning.

And the unglamorous bit that matters most: **it happens in time that did not previously exist.** You did not read faster — you found hours.

Answer 4: how do you use AI day to day

ABOUT 85 SECONDS

Thanks for watching. The thread through how I use it is that I let it do the sorting and I keep the deciding.

So a small everyday one — I was getting a huge volume of email from university and everywhere else, and the problem wasn't reading them, it was working out which ones actually needed me. So I connected Claude to my inbox and had it colour-code by type: needs action from me, just CC'd, has a deadline attached. It doesn't reply to anything, it just sorts — but it means I can plan the day round the ones that are actually mine.

And for learning, which is most of how I use it — I'll use Claude Code to pull together material on a topic, then put that into NotebookLM and get an audio version I can listen to at the gym. That gives me a baseline. Then the important bit: I go back and read properly on the parts I didn't follow or wanted more on. So the audio is really a way of finding out what I don't know, and the reading is where the actual learning happens.

Two things I'd say about that. One, it hasn't made me read faster, it's just meant the reading happens in time that wasn't there before. And two, I've got a rule that I won't quote anything I haven't read in the original, because summaries drop the caveats and the caveats are usually the interesting bit.

PART V

Allocation and delivery

9 Which answer for which question

If the question is	Give
"A time you used AI to solve a problem"	**Answer 1** — Pictet, in full.
"How do you make sure output is accurate" / "what are the risks"	**Answer 2.** It reuses Pictet as evidence but is *not* a retelling, so it survives even if answer 1 came first.
"How will AI change asset management / this role"	**Answer 3.** One clause of Pictet, no more.
"How do you use AI day to day / to learn"	**Answer 4.** Gmail and the research workflow.

THE REPETITION RISK — READ THIS

Pictet appears in three of the four answers. **That is fine, because it appears differently each time** — as the story, as evidence for a principle, as one clause of illustration. But if you tell the full DDQ story twice, the second one sounds like you have one anecdote.

The rule: full story once. After that, a single sentence maximum. "I saw a version of that at Pictet" is enough.

10 Tonight

1. **Settle whose idea the DDQ project was.** You said "as an intern, I began talking to my team, and I realised" — if that is accurate, you have genuine initiative and should say it plainly. If the brief came to you, say that and own the design instead. Either is strong; a mismatch at Superday is not.
2. **Decide how far it actually got.** You said you escalated for company approval. Was it approved? In use? Handed over? Say exactly, and do not round up.
3. **Memorise four things per answer, not sentences.** For answer 1: 400–500 questions, five emails per misroute, routing not writing, low-consequence by design.
4. **Film each one eight to ten times**, changing the words, keeping the beats fixed.

> **THE FOUR SENTENCES THAT CARRY ALL OF THIS**
>
> **1.** "The bottleneck wasn't writing the answers, it was routing them."
> **2.** "I put it on the high-frequency, low-consequence bit and kept people on the other one."
> **3.** "When everyone can produce a polished response, the polished response stops being what distinguishes you."
> **4.** "It does the sorting, I do the deciding."
>
> If you land those four, you will be well clear of a field that mostly says AI saves them time.

What is mine and what is yours. Every fact is yours — the Excel tracker, the 400 to 500 questions, the five-email ping-pong, the dashboard and the reassignment design, the sandbox, the escalation, the Gmail colour-coding and the research workflow. What I have added is the framing: that you diagnosed a routing problem rather than a writing problem, that you designed yourself out of the loop, that you chose a low-consequence use case, the risk ranking, the two verification principles, the three-link chain on what AI changes and the conclusion that it raises rather than lowers the premium on the relationship. The chain in section 7 is my reading, not a published view — defend it as a position, not a fact. The arithmetic in section 2 uses your own numbers and an assumption you should flag when you say it. Everything in square brackets is yours to supply and should not be said unless true. One disclosure repeated from doc 47: I am Claude, made by Anthropic, so weigh my advice about naming AI tools accordingly — name them for specificity, not because any tool is impressive.