Accuracy, Reliability and Risk

"How do you make sure AI output is accurate?" and "What are the risks?" — the two questions where the obvious answer is the losing one, and the mechanism that replaces it.

Doc 54 · Builds on doc 50 (the four AI answers) and doc 53 (the industry view) · Prepared 6 November 2026

IF YOU READ NOTHING ELSE BEFORE YOU RECORD

1. The textbook answer — "I always double-check it, I use it as a starting point not an end point" — is this question's version of corporate slop. Unfalsifiable, unmemorable, everyone says it.

2. The answer that works is a **mechanism, not a habit**: *"The real discipline isn't checking harder. It's only using it where checking is cheap."*

3. And the reflection that lands is the pragmatic trade-off — *"I didn't make the model more accurate. I made being wrong cheap to fix."* That is what your DDQ dashboard actually did.

WHAT IS IN HERE

1. **Part I** — Why the obvious answer fails, and how your own framework fixes it.
2. **Part II** — Accuracy as a design problem: the four-part mechanism, with the grid.
3. **Part III** — The risks, sorted by who actually pays for them.
4. **Part IV** — The commercial turn: your risk list is now the client's questionnaire.
5. **Part V** — Two scripted answers, a merged one, and the six lines you asked for.
6. **Part VI** — What not to say, follow-ups, and the honesty guardrail.

PART I

Why the obvious answer fails

1 This question's version of corporate slop

The framework you sent makes a sharp point about behavioural questions: nobody wins by claiming *"I sat them down, motivated them, and they became our highest performer."* Senior people know human nature does not work like that, so the claim reads as fabricated.

The same trap exists here, and almost nobody spots it. Here are the four things most candidates will say on camera today:

The slop	Why it fails
"I always double-check the output."	Unfalsifiable and almost certainly untrue. Nobody double-checks everything. Claiming it tells the reviewer you have not thought about the cost of checking.
"I use it as a starting point, never an end point."	A slogan. It describes an attitude, not a process. The follow-up "how, specifically?" destroys it.
"I always verify the sources."	Only meaningful if you say what you do when there is no source to verify — which is the usual case.
"You have to use human judgment alongside it."	True of literally everything. Says nothing.

All four describe **a habit you are claiming to have**. A habit cannot be checked, so it carries no information. What carries information is **a mechanism you built** — because a mechanism can be described, questioned, and found to exist.

2 The move that works, translated from your own framework

Your framework says the differentiator in a behavioural answer is the sophisticated reflection: *admit the real constraint, defend the pragmatic trade-off, anchor it to an outcome.* Sam Leong's version was "taking over someone's workload isn't the textbook solution, but it was the only realistic way to protect against a plagiarism fail."

Translate that move to this question and you get:

THE STRUCTURAL EQUIVALENT

Textbook claim: "I verify everything the model produces."

Pragmatic trade-off: "I couldn't verify everything — so instead of pretending to, I restricted what I used it for. I pointed it at work where being wrong was cheap to catch, and kept a person on the work where it wasn't."

That is the same shape of answer. It admits a constraint most candidates pretend does not exist, defends a decision made because of that constraint, and it is **true**, which means it survives every follow-up.

WHY THIS IS UNUSUALLY STRONG FOR YOU SPECIFICALLY

Most candidates have to invent a verification process, because they were using AI for essays. **You actually built one.** The reassignment button on your dashboard is a verification mechanism — you just did not call it that.

You did not make the routing more accurate. You made a wrong routing cost one click instead of five emails. That is the entire answer, and you already have it.

PART II

Accuracy as a design problem

3 The four-part mechanism

One — choose tasks where checking is cheaper than doing

This is the foundation and nearly everyone misses it. The question is never "how carefully did you check?" It is "how expensive was checking in the first place?"

- **Cheap to check:** routing a question to the right subject matter expert. The person who receives it knows in two seconds whether it belongs to them.
- **Expensive to check:** summarising a two-hundred-page document. Verifying the summary costs roughly as much as reading the document — so in practice nobody does, and the "verification" is theatre.

So the first discipline is a selection decision, made before you touch the tool.

Two — build the correction path into the thing itself

Do not put the safety net outside the system, where it depends on somebody remembering. Put it inside, where using the system *is* the check.

Your dashboard did this. A question got routed; if it was wrong, the recipient reassigned it in one click. The error was surfaced by the person best placed to spot it, at the moment they were already looking, and fixing it was trivial. No separate review step, no reliance on diligence.

Three — a named human owns anything that leaves the building

Internal draft and client-facing output are different categories and deserve different bars. Something that goes to a client is a representation the firm is making. Somebody signs it. That does not change because a model produced the first draft.

Four — know the category of things it cannot know

The sharpest version of this: **a model cannot know anything that was never written down.** Not this quarter's actual position. Not what a client said on a call last week. Not what your colleague decided yesterday afternoon.

So the failure mode people describe — "it makes things up" — is slightly wrong. The real failure mode is that **it confidently answers a question it had no basis to answer**, because it has no mechanism for saying "I was never told."

4 The grid — and where your Pictet work sits on it

Two things determine whether a task is safe to hand over: how much it costs if the output is wrong, and how much it costs to find out.

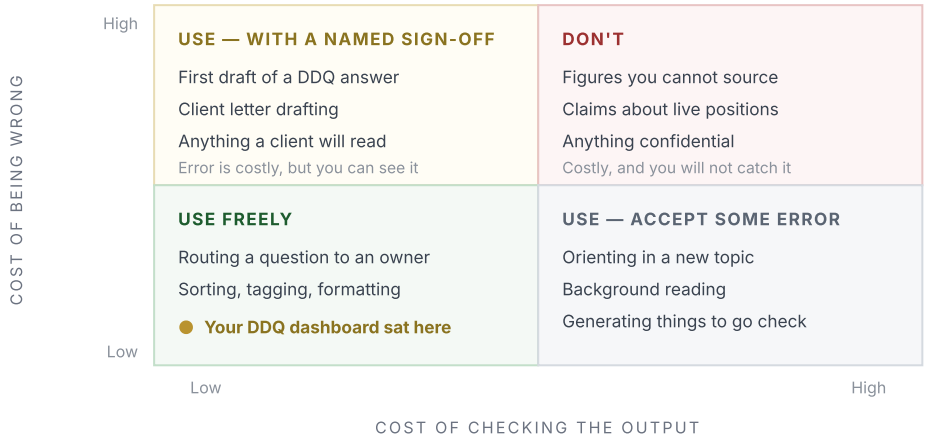


Figure 1 — Accuracy is mostly a task-selection decision, made before you open the tool.

Note what the grid says about the bottom-left box. It is not safe because the model is good there. It is safe because **being wrong there costs nothing**. That distinction is the whole insight, and it is worth saying explicitly on camera.

> **THE HONESTY GUARDRAIL — READ THIS TWICE**
>
> Only claim the checking you actually did. If you did not review every routing suggestion, do not say you did. The pragmatic version is stronger anyway: *"I didn't check every one — I didn't need to, because the person receiving it would."*
>
> And Saad's rules still hold, in this answer as in every other: **no Pictet client names**, nothing confidential, no claiming the project was your idea if your manager suggested it. If the tool was firm-approved or you only used non-confidential data, say so — but only if it is true.

PART III

The risks, sorted by who pays

5 A better cut than a flat list

Doc 50 ranked the risks. Here is a more useful way to organise the same material for an interview, because it shows you can think about a risk from somebody else's balance sheet rather than just your own.

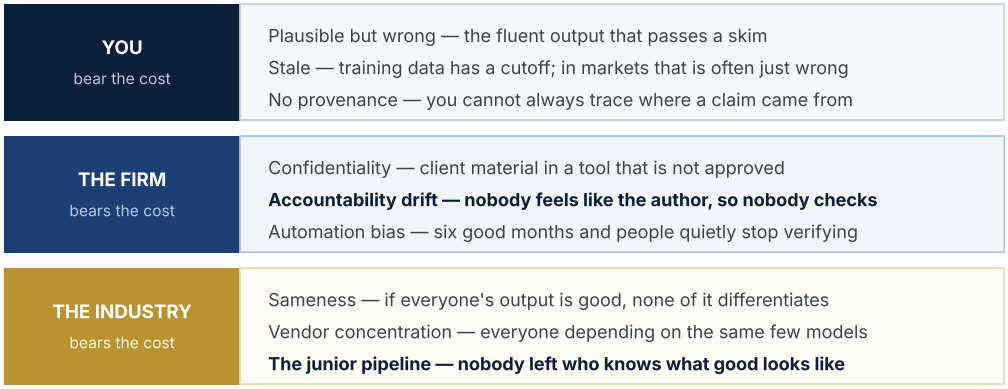


Figure 2 — Same risks, organised by whose problem they become. The two in bold are the ones worth leading with.

6 The two risks worth leading with

Accountability drift

This is the best one, and very few candidates will have it.

When a person writes something, they feel like the author. They have skin in it. When a machine drafts it, nobody does — and the output *looks finished*, which makes it feel like it has already been through a process it has not been through.

> The risk is not that the model is wrong. It is that **a draft arrives looking like a final document**, so it receives less scrutiny than a human first draft would have, precisely because it is better formatted.

That is an organisational risk rather than a technical one, which is what makes it interesting. And it connects straight to the regulatory position from doc 53: regulators are not writing new AI rules, they are saying existing accountability rules apply, and the FCA expects a firm to be able to name the senior manager responsible for each significant AI use. The regulatory answer to AI is *accountability* — the same conclusion, arrived at from the other direction.

The junior pipeline

The slower one, and it earns credit for being uncomfortable.

Producing the first draft is how junior people learn what a good one looks like. If the machine always produces it, that learning does not happen — and in five or ten years there is nobody who can tell whether the machine's output is any good.

This ties to the EY finding in doc 53: firms are reducing junior headcount while investing in senior hires, especially in sales. Which is a sensible short-term decision that may have an awkward long-term consequence. Saying that shows you can hold two time horizons at once.

USE ONE OF THESE, NOT BOTH

Each takes about twenty seconds to deliver properly. Pick accountability drift as the default — it is tighter and it links to the regulatory point. Keep the junior pipeline in reserve for a follow-up, or if the question is specifically about the industry rather than about you.

PART IV

The commercial turn

7 Your risk list is now the client's questionnaire

Here is the move that converts this from an IT answer into a commercial one, and it is the single most valuable thing in this document.

The risks above are not abstract. **They are exactly what institutional clients now ask asset managers in writing.**

- The Investment Company Institute's *Investment Funds Distributor Due Diligence Questionnaire* (August 2026) has an AI section asking whether AI or machine-learning tools are used in distribution activities, whether they were built in-house or by a third party, and whether policies govern their use. MEDIUM
- Robeco's published question set asks where AI is used, which decisions it influences, who owns AI governance, what the failure protocol is, and whether the manager can reconstruct the data, assumptions, outputs and human approvals behind a material AI-assisted process. MEDIUM
- Mercer adds vendor concentration — whether a manager understands its downstream AI dependencies. MEDIUM

WHY THIS IS THE STRONGEST THING YOU CAN SAY ON THIS QUESTION

Every other candidate answers "what are the risks of AI" as though it were a personal opinion question. You can answer it as a **client question** — because that is now literally what it is.

One sentence does it: *"What's interesting is that this has stopped being an internal governance question — clients now ask it directly in due diligence questionnaires, including whether you can reconstruct the human approvals behind an AI-assisted process."*

That sentence tells the reviewer three things at once: you know what a DDQ is, you know what is currently in them, and you understand that risk management is a *sales* issue in this business, not just a compliance one.

PART V

8 Answer A — how do you make sure it is accurate and reliable?

ABOUT 85 SECONDS — KEEP THIS ONE TIGHTER THAN YOUR STORY ANSWERS

"Thanks for watching. I'd start by saying the risk isn't an output that's obviously wrong — those get caught. It's the one that's fluent, well-formatted and subtly incorrect, because that's the one that passes a skim and gets used.

So the thing I've landed on is that the real discipline isn't checking harder. It's only using it where checking is cheap. When I worked on the due diligence questionnaire process on my placement, I deliberately pointed it at routing questions to the right subject matter expert rather than writing the answers. If a routing is wrong, the person who receives it knows in about two seconds, and reassigning takes one click. If a client-facing answer is wrong, that's a representation the firm is making and you might not find out for years — so that part kept a person on it.

And I'd be honest that I didn't verify every suggestion. I didn't need to, because the correction was built into how the thing worked rather than sitting outside it as a review step somebody had to remember. I didn't make the model more accurate. I made being wrong cheap to fix.

The other rule I keep for my own reading is that I don't quote anything I haven't read in the original. Summaries drop the caveats, and the caveats are usually where the interesting part is."

Why this clears the bar. It names a specific failure mode, gives a mechanism rather than a habit, admits what was not checked, and anchors to a design decision you actually made. The last paragraph is a second, independent discipline, which shows this is a way of thinking and not one anecdote.

9 Answer B — what are the risks of using AI?

ABOUT 85 SECONDS

"Thanks for watching. I'd split it into three, by who actually bears the cost.

The one I bear is plausible-but-wrong — fluent, confident output that's subtly incorrect. Obvious errors are safe; it's the ninety-five-percent-right answer that's dangerous, because it passes a skim.

The one the firm bears is more interesting, and I'd call it accountability drift. When a person writes something they feel like the author. When a machine drafts it, nobody does — and because it arrives looking finished, it can get less scrutiny than a human first draft would have. That's an organisational risk rather than a technical one. And it's notable that the regulators have landed in the same place: they're largely not writing new AI rules, they're saying existing accountability rules apply, and expecting firms to name who's responsible for each significant use.

Then there's confidentiality, which is the non-negotiable one — client material only ever goes into approved, contained tools.

And what strikes me is that this has stopped being an internal governance question. Clients now ask it directly in due diligence questionnaires — including whether a manager can reconstruct the data and the human approvals behind an AI-assisted process. So getting this right isn't just risk management. It's part of how you win the mandate."

10 If they are asked together — the merged version

ABOUT 110 SECONDS, IF THE PROMPT BUNDLES BOTH

"Thanks for watching. I'll cover the risk I think is underrated, what I actually did about it, and why it's become a commercial question rather than just a technical one.

The risk isn't the output that's obviously wrong — those get caught. It's the fluent, well-formatted one that's subtly incorrect, because it passes a skim. And underneath that sits a quieter one: when a machine drafts something it arrives looking finished, so nobody feels like the author and it can get less scrutiny than a human first draft would have.

What I took from that is that the real discipline isn't checking harder — it's only using it where checking is cheap. On my placement, working on the due diligence questionnaire process, I pointed it at routing questions to the right expert rather than writing the answers. A wrong routing is caught in two seconds by the person who gets it, and fixed in one click. A wrong client-facing answer is a representation the firm is making, so that kept a person on it. I didn't verify every suggestion and I wouldn't claim I did — I made being wrong cheap to fix rather than making the model more accurate.

And the reason I think this matters commercially is that clients now ask about it in writing. Due diligence questionnaires have started including AI sections — whether tools are built in-house or third-party, who owns the governance, whether you can reconstruct the human approvals behind an AI-assisted process. So this has become part of winning the mandate, not just part of managing the risk."

11 The six lines — what you asked for

These are the droppable sentences. Memorise the sense of them, not the words. Each one does work no other sentence in your answer does.

> *"The dangerous output isn't the one that's wrong. It's the one that's ninety-five percent right, because that's the one that passes a skim."*
>
> **WHY IT LANDS** Most candidates describe AI errors as obvious. Saying the opposite shows you have actually used it on something that mattered.

> *"The real discipline isn't checking harder. It's only using it where checking is cheap."*
>
> **WHY IT LANDS** Converts a vague virtue into a selection rule. This is the line to build the whole answer around.

> *"I didn't make the model more accurate. I made being wrong cheap to fix."*
>
> **WHY IT LANDS** Describes a design decision, not an intention — and it is a true description of what you built. This is your pragmatic trade-off.

> *"It can't know anything that was never written down — so the failure mode isn't invention, it's confidently answering a question it had no basis to answer."*
>
> **WHY IT LANDS** A precise correction to the usual "it hallucinates." Precision reads as experience.

> *"When a machine drafts it, nobody feels like the author — so it can go out with less scrutiny than a human first draft, precisely because it looks more finished."*
>
> **WHY IT LANDS** An organisational risk rather than a technical one. Almost nobody will say this, and it is the one that sounds like someone who has worked in a firm.

> *"This has stopped being an internal governance question — clients now ask it in due diligence questionnaires."*
>
> **WHY IT LANDS** Turns a risk answer into a commercial answer, and quietly proves you know what a DDQ is and what is currently in one.

USE THREE, NOT SIX

Six strong lines in ninety seconds is a recital, and it will sound like one. **Pick lines 2, 3 and 6** as your spine — the rule, the trade-off, the commercial turn — and hold the rest for follow-ups.

PART VI

12 What not to say

Don't say	Say instead
"I always double-check everything."	"I didn't check every one — I didn't need to, because the correction was built in."
"The main risk is hallucination."	Name the specific failure: fluent, plausible, subtly wrong. Then go to accountability drift.
"It's just a tool, you still need judgment."	Nothing. Cut it. It is true of everything and signals nothing.
Any Pictet client name, or anything confidential	"A client questionnaire." Saad was explicit about this.
Claiming the DDQ project was your idea, if it was not	"One of my tasks was…" or "a senior analyst and I…" — which is also the truth, and reads better.
Your AI workflow for preparing for this interview	Leave it out entirely. In every answer.
"AI will replace a lot of these jobs."	The doc 53 version: the document-producing half is automated, the accountable half is being invested in.
Over-engineering it into a governance lecture	One mechanism, one honest admission, one commercial point. Then stop.

13 Follow-ups you might get

If they ask	One-line answer
"Give me an example of when it got something wrong."	Use a real routing miss if you have one — a question that went to the wrong desk and came back. The point is that the system caught it, not that you did.
"So did you check all of it or not?"	"No — and that was deliberate. Checking every suggestion would have cost more than the process saved. The design assumed some would be wrong." Confident, not apologetic.
"What would you do differently?"	Honest option: log the corrections. If you can see which questions get reassigned most, you learn where the process itself is ambiguous — which is a better finding than the routing.
"Isn't confidentiality the biggest risk?"	"It's the most serious, but it's also the most solved — firms have approved tools and clear rules. The one that's less solved is who's accountable for an output nobody feels they authored."
"How would you explain this risk to a client?"	"I'd tell them what we use it for, what we don't, and who signs the thing that reaches them. The first two are policy. The third is the one they actually care about."
"Do you think AI makes the industry safer or riskier?"	Both, honestly — fewer manual errors, more correlated ones. If everyone uses similar tools, mistakes stop being independent.

14 Before you record

1. **Decide which AI question gets Pictet.** Doc 47 flagged this. If the "how have you used AI" question comes first and takes the DDQ story in full, this answer leads with the *principle* and references the project in one clause — not the full story twice.
2. **Say the honest admission out loud once.** "I didn't verify every suggestion" feels risky until you have heard yourself say it. It is the strongest sentence in the answer.
3. **Keep it to three lines from section 11.** Two, the rule. Three, the trade-off. Six, the commercial turn.
4. **Do not lecture.** This answer should be shorter than your story answers. Eighty-five seconds is right; two minutes of governance commentary is not.
5. **End on the client, not the technology.** Every one of your answers should land back on the person whose money it is.

Sourcing and confidence. MEDIUM: the ICI *Investment Funds Distributor Due Diligence Questionnaire* (August 2026) AI section and its contents; the Robeco, Aon and Mercer investor question sets on AI governance, including reconstructability of data, assumptions, outputs and human approvals, and Mercer's vendor concentration point; the FCA position that existing accountability rules apply to model-made decisions and that firms should be able to name the responsible senior manager; the EY 2026 finding on junior headcount versus senior sales hiring. All are carried over from docs 53 and 50 and were sourced from secondary reporting and published practitioner material rather than primary regulatory texts — re-verify before quoting a specific document name on camera, and prefer the general

formulation ("questionnaires have started including AI sections") over naming the ICI document. **MINE**: the slop table in section 1, the translation of the STAR-R pragmatic trade-off to this question in section 2, the four-part mechanism, the cost-of-checking grid in figure 1, the who-bears-the-cost cut in figure 2, the accountability-drift framing, the junior-pipeline argument, all six lines in section 11, and both scripted answers. The Pictet details are yours — I have described only what you told me, and you should check every factual claim about the project before you say it.