The Real Questions

All six prompts mapped, the allocation problem solved, and the two scenario questions taken apart — including why they are designed so that you cannot solve them, which is the thing that stops the panic.

Doc 56 · Goldman Sachs Asset Management HireVue, actual question set · Builds on doc 46's four-beat method · Prepared 6 November 2026

THE THING TO READ BEFORE ANYTHING ELSE

Both scenarios are built so that you cannot solve them. The manager is unreachable by construction. The department is unreachable by construction. The deadline is tonight by construction. The team has already agreed by construction.

If you go in thinking the task is to *solve* it, you will panic — because it cannot be solved, and part of you knows that. The task is to **handle** it. Those are different jobs, and only one of them is possible.

Once you see that, the whole thing calms down. You are not being asked for the answer. You are being asked what you would do when there isn't one.

WHAT IS IN HERE

1. **Part I** — All six questions mapped, and the story allocation sorted.
2. **Part II** — Why scenario questions exist, and what is actually being scored.
3. **Part III** — The protocol: what to do with your 45 seconds of prep time.
4. **Part IV** — Question 2, the data discrepancy, in full.
5. **Part V** — Question 3, the AI section, in full.
6. **Part VI** — Question 5, which you have not prepped and may not have noticed.

PART I

1 What each question is actually testing

#	The question	What it is testing	Your material
1	Walk me through your resume	Can you structure, prioritise and not ramble. Floor question — easy to fail, impossible to win	Doc 28 (playbook), doc 17 (CV)
2	Data discrepancy before a client meeting	Judgment under pressure. Do you ship a known problem, freeze, or escalate well	**This doc, Part IV**
3	Team wants to submit unreviewed AI work	Will you hold a line against a tired consensus — without being insufferable about it	**This doc, Part V**
4	AI to support a teammate or solve a team challenge	Real use, not theory. Plus whether you make others better or just yourself faster	Doc 50, doc 54. **Pictet DDQ**
5	Extended project — staying engaged, meeting deadlines	Stamina and self-management over months, not a sprint	**See Part VI — you have not prepped this**
6	Why asset management, which area, what skills	Commercial motivation. The one you can genuinely win on	Docs 42, 49, 13, 48

2 The allocation problem, and the good news in it

Look at what is **not** on this list: no teamwork question, no failure question, no weakness question, no conflict question, no "above and beyond", and no market news question.

That means a lot of the work you did is not needed today — which is fine, because it was never wasted (it is first-round interview material) — and more importantly it means **your stories are no longer competing with each other.**

Story	Was prepped for	Now
Pictet DDQ	AI question and challenge question — a clash	Goes to **Q4**, uncontested. The challenge question isn't asked
House hunt	Teamwork	**Free.** Teamwork isn't asked — so it goes to **Q5**, and it fits better than you'd think
Society event, Dubai payment, slide deck, conflict	Failure, above and beyond, weakness, conflict	Not needed. Hold them for the next round

THE ALLOCATION THAT FALLS OUT OF THIS

Q4 → Pictet DDQ. It is an AI story *and* a team-challenge story, which is exactly what the question asks for. Lead on how it helped the team, not how clever the build was.

Q5 → the house hunt. Months long, a hard external deadline, and everybody else disengaged while you kept it moving. That is precisely what "stayed engaged and ensured deadlines were met" is asking about.

No story does double duty. That is the first time that has been true in this whole process.

3 One trap in question 6 before we leave the map

Question 6 is **three questions in one**: why asset management, which area of the business, and what skills make you successful. Candidates routinely answer the first, forget the second, and give a generic version of the third.

GIVE EACH PART AN EXPLICIT SIGNPOST

"I'll take those in turn." Then: why asset management (docs 42 and 49 — duration of fees, whose money is at risk). Which area (Client Solutions, and say why specifically — doc 48 and doc 13). What skills (pick **two**, each with a one-line proof, not a list of five adjectives).

If you run short on time, sacrifice part one. They can infer why you like asset management. They cannot infer which desk you want.

PART II

Why scenario questions exist

4 Why firms ask these, and why they weight them heavily

Every other question in this set can be pre-written. You can draft a resume walkthrough, rehearse a why-asset-management answer, polish an AI story for a fortnight. The firm knows that.

The scenario questions are **the only part of the interview you cannot script in advance**. Which makes them the highest-signal part of the whole exercise, and it is why they are in there.

There is a second reason, and it is more interesting. Behavioural questions tell the reviewer what you *did* once, in a situation you chose to describe and have had time to make flattering. Scenario questions tell them **how you reason when you do not know the answer** — which is a much better predictor of what you will be like as a first-year analyst, because being a first-year analyst consists almost entirely of not knowing the answer.

5 What is actually being scored

Not whether you got it "right." There are five things the reviewer is listening for, and you can hit all five without resolving anything.

What they listen for	What it tells them
Do you name the stake first?	Whether you think about consequences before actions. This is the single biggest differentiator and it is the opening beat.
Do you act, or freeze?	A candidate who says "there's nothing I could do" has failed. There is always something.
Do you act *alone*?	Equally bad the other way. A junior who resolves things unilaterally is a risk, not an asset.
Is your response proportionate?	Do you scale the reaction to how serious it is — or treat every problem as a crisis and every crisis as a formality?
Do you handle the people as well as the task?	Almost every scenario has a procedural half and a human half. Most candidates do one and forget the other.

6 The reframe that stops the panic

You said you panic on these. Here is the reason, and the fix.

You panic because you believe you are being asked to produce **the correct answer**, in sixty seconds, to a problem you have never seen. That would be frightening, and it would also be impossible.

> But there is no correct answer. There is **a correct set of considerations**, and you are scored on whether you surface them in the right order. That is a completely different task — and it is one you can prepare for, because the considerations are always the same four.

Which means you are not inventing an answer under pressure. You are **filling in four slots**. The content changes with the question. The slots never do.

PART III

The protocol

7 What to do with the 45 seconds of prep time

This is the practical anti-panic tool. When the scenario appears, do not try to think of an answer. Fill the four slots in order, spending roughly ten seconds on each.

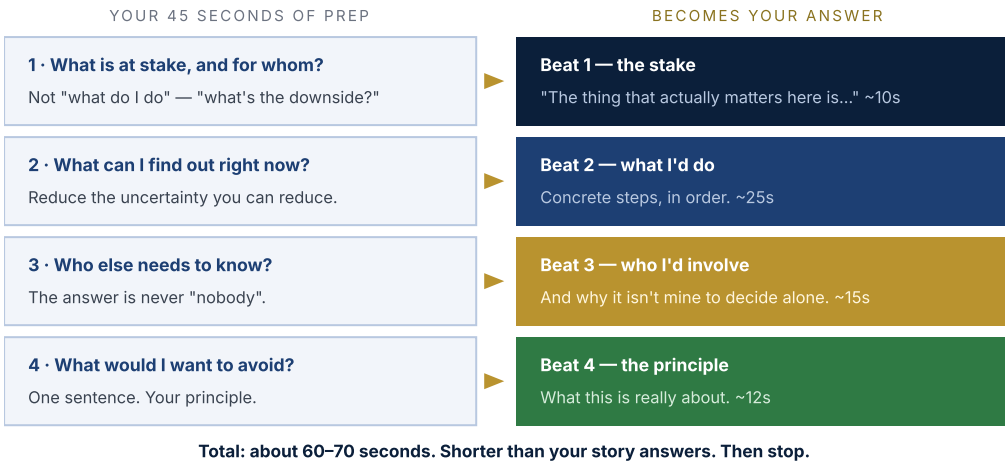


Figure 1 — The prep time and the answer have the same four slots. You are not composing; you are filling in.

WHY BEAT 1 DOES MOST OF THE WORK

Starting with the stake rather than the actions is what makes an answer sound like judgment instead of rule-following. It also solves your panic directly: **you always have a first sentence**, even on a scenario you have never imagined, because naming what is at risk requires no cleverness at all.

If you freeze, say the stake out loud. The rest follows from it.

8 Four rules that cover almost every failure

1. **Never act unilaterally on something that is not yours to decide** — but equally, "I'd ask my manager" cannot be the *whole* answer. Show what you would do *and* who you would involve.
2. **Do not be absolutist.** "I would never…" sounds naive. Real judgment has conditions attached.
3. **Say what you do not know.** "I don't know what the firm's process is, so I'd find out" is a strong answer, not a weak one. It is also true, which helps.
4. **Keep it short.** These are pass-or-fail. Answer well, then stop. Elaborating is where people talk themselves into the wrong answer.

PART IV

Question 2 — the data discrepancy

THE PROMPT

Your manager is in a meeting until they leave for the day and is unreachable. You're finalising a presentation they need for a client meeting later that day. During a final check, you discover a discrepancy in data provided by another department, and you can't reach that department before the meeting begins. What do you do?

9 What is actually at stake

Goldman's entire franchise rests on one proposition: **when the firm puts a number in front of a client, the number is right.** Not "usually right." Right.

So the hierarchy here is not subtle, and saying it out loud is most of beat 1:

> A presentation that is **late** is an inconvenience. A presentation that contains a known error, presented by someone who did not know it was there, is a **client trust problem and potentially a regulatory one**. Those are not close.

And notice who is most exposed: not you. **Your manager**, who will be standing in front of a client asserting a figure they do not know is in question. Saying that shows you understand the situation from somebody else's position, which very few candidates do.

10 The three traps in this question

Trap one: taking "unreachable" at face value

This is the big one, and almost everybody falls into it.

The prompt says your *manager* is unreachable. It does not say **everyone** is. In a real firm there is always a second senior person: another member of the team, a desk head, your manager's manager, the assistant who can pass a note into a meeting room.

Same for the other department. "I couldn't reach the department" usually means you emailed one person. Departments have shared inboxes, team leads, and people who sit two desks away from the person who sent the data.

> **THE SENTENCE THAT SETS YOU APART ON THIS QUESTION**
>
> *"My manager being in a meeting doesn't mean nobody senior is available — so I'd find the next person who could make that call, rather than treating one closed door as the end of the options."*
>
> Almost no candidate says this. It reads as someone who has worked somewhere real.

Trap two: believing you have to resolve the discrepancy

You probably cannot. The data is not yours, the people who own it are unavailable, and the meeting is today. If you define success as "fix the number," you will give a flustered answer.

You do not have to fix it. You have to make sure **nobody walks into that room and asserts something they cannot stand behind.** That is entirely within your control, and it is the actual deliverable.

Trap three: treating every discrepancy as equally serious

It is not. Before you escalate anything, work out whether it **changes the conclusion**. A rounding difference in a footnote and a wrong headline figure the client will act on are different problems and deserve different responses. Scaling your reaction to materiality is what makes you sound senior rather than anxious.

11 What you can and cannot control

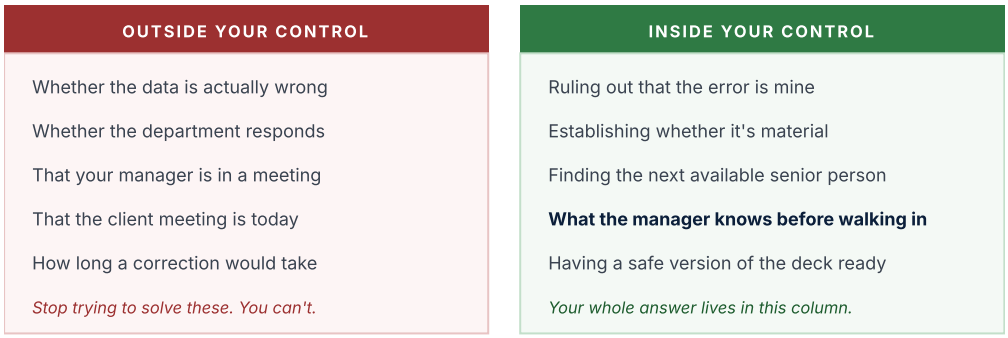


Figure 2 — The scenario is engineered so the left column is impossible. The answer is entirely on the right.

12 The full answer

ABOUT 75 SECONDS — RESIST THE URGE TO MAKE IT LONGER

"The thing that actually matters here is that a presentation being late is an inconvenience, but a presentation with a known error in it, presented by someone who doesn't know it's there, is a client trust problem. And the person most exposed is my manager, because they're the one who'd be asserting it in the room.

So first, before escalating anything, I'd spend a few minutes on what I can actually establish. Is the error mine — a transcription or a formula on my side? Has the number changed between drafts? And crucially, is it material — does it change the conclusion the client would act on, or is it cosmetic? That determines how hard I push on everything else.

Then I'd widen the search for a person. My manager being in a meeting doesn't mean nobody senior is available — I'd try another member of the team or whoever's covering, and on the other department I'd go to the team or a shared inbox rather than the one person who sent it.

And while that's running, I'd make the deck safe. I'd have two versions ready — one with the figure flagged and caveated, one with it removed — plus a short written note: here's what I found, here's what I checked, here's what I couldn't verify. Even if I only get to my manager as they're leaving, they walk in knowing.

I think the principle is that I probably couldn't fix the number in the time available. What I could control was that nobody asserted something we couldn't stand behind."

ONE PHRASING CHECK

That closing line has the same shape as your Dubai reflection ("I didn't change the outcome, I changed the not-knowing"). That is fine — it is how you think, and consistency is good. But do not use *identical wording* in two answers in the same recording, because the reviewer watches them back to back and it will sound like a formula.

PART V

Question 3 — the AI section

THE PROMPT

Imagine you're finishing a major group project, and the night before it's due, a teammate suggests using AI to generate the entire final section and submitting it without review or disclosure. The rest of the team agrees because everyone is exhausted. How might you handle the situation, and what alternatives might you suggest?

13 There are two problems here, and candidates merge them

This is the unlock for the whole question. Almost everyone treats this as one thing — "using AI to cheat" — and then argues against AI, which is both wrong and unpersuasive.

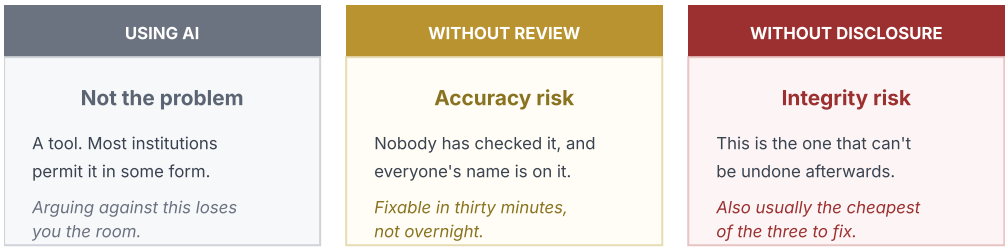


Figure 3 — Three separate things in one sentence. Splitting them is what makes the answer workable.

Separating these is the whole answer. **The tool is not the issue. "Without review" and "without disclosure" are two different issues** — and both can be fixed in under an hour, which is what makes your alternative actually acceptable to an exhausted team.

14 The four wrong answers

The wrong answer	Why it fails
Go along with it	Disqualifying. There is no version of this that passes.
The moral lecture "I'd explain that academic integrity is extremely important"	Correct and completely useless. It addresses the ethics while ignoring the actual problem, which is that four people are exhausted and it is due tomorrow. You lose the argument because you have offered them nothing.
Straight to the professor	Escalating before attempting to solve it reads as someone who *creates* problems rather than resolves them. Escalation is the backstop, not the opening move.
"I'd have a conversation with the team"	Says nothing. The question explicitly asks what alternatives you would suggest — so suggest some.

15 Address the real problem, not the stated one

Here is the thing most candidates never see. **The team is not proposing this because they are dishonest. They are proposing it because they are shattered and out of road.**

So if you argue ethics, you lose — not because you are wrong, but because they are too tired to be moved by principles. You win by making the honest option *cheaper than the dishonest one*.

And there is one more point worth making, which comes straight from the risk material in doc 54:

> **THE ARGUMENT THAT ACTUALLY PERSUADES TIRED PEOPLE**
>
> *"The risk isn't that the AI section is bad. It's that it's plausible — nobody skim-reading an unreviewed section spots the one fabricated citation. Which is why 'unreviewed' is the dangerous word in that sentence, not 'AI'."*
>
> That reframes it from a lecture into a warning about **their** exposure. Exhausted people will not accept a moral argument at midnight. They will accept a risk argument.

16 The four alternatives — because the question asks for them

Alternative	What it costs
Use it, but review it	Drafting with AI and then checking is a completely different act from submitting unchecked. Split across four people it is fifteen minutes each, not an all-nighter.
Disclose it	If the institution permits AI with a declaration, declare it. This is usually the cheapest fix and it removes the part that cannot be undone.
Shorten the section	A shorter section the team can stand behind beats a longer one it cannot. Markers do not reward length.
Ask for an extension	Frequently available, rarely requested, and the worst outcome is being told no. Costs one email.

17 The full answer

ABOUT 85 SECONDS

"I think there are actually three things in that sentence and they're not the same. Using AI isn't the problem — most institutions permit it in some form. Submitting it without review is an accuracy problem, and submitting it without disclosure is an integrity problem. The second one is the one you can't undo afterwards.

So I wouldn't open by lecturing anybody, because everyone's exhausted and a moral argument at midnight doesn't land. I'd make the risk argument instead: the danger isn't that the section is bad, it's that it's plausible. Nobody skim-reading an unreviewed section catches the one fabricated citation — and all four of our names are on it.

Then I'd put alternatives on the table, because that's the part that actually changes anyone's mind. Use it, but review it — split across the team that's fifteen minutes each, not an all-nighter. Disclose it, if the policy allows AI with a declaration, which removes the unrecoverable bit for almost no cost. Shorten the section if we have to — a shorter piece we can stand behind beats a longer one we can't. And ask for an extension, because the worst answer is no.

If the team still wanted to submit it undisclosed, I wouldn't put my name to it, and I'd say so plainly rather than going round them. But I'd want that to be the last resort, not the opening move — it's a worse outcome for everyone, including me.

The general lesson I'd take is that the tool was never the issue. It's whether somebody owns the output."

TONE — THIS MATTERS AS MUCH AS THE CONTENT

Deliver this **practically and slightly sympathetically**, not righteously. "Everyone's exhausted and I get it" should be audible in how you say it. A candidate who sounds pleased to be the only honest person in the room scores badly, even when the content is right.

And note the last resort is phrased as a last resort, with a reason. "I'd refuse immediately and report them" is a worse answer than it feels like.

PART VI

Question 5 — the one you have not prepped

THE PROMPT

Describe a project that required your attention for an extended period. How did you stay engaged and ensure deadlines were met?

18 Why this one is easy to under-prepare

You have been focused on the scenarios, which is reasonable. But this question is the one with **no prepared story attached to it** — Pictet is going to Q4, and everything else you built was for questions that are not being asked.

It is also a two-part question, and the parts are different: *how did you stay engaged* is about you over time, and *how did you ensure deadlines were met* is about the mechanism. Answer both.

19 Use the house hunt, and here is why it fits better than it looks

Six people, months of effort, a hard external deadline (good properties go), you in London and the property in Norwich, and most of the group not responding. You absorbed the work rather than trying to split it evenly, and reduced everything to yes/no decisions so people could actually answer.

What the question asks	What the house hunt gives you
An extended period	Months, not a sprint. Genuinely sustained.
How you stayed engaged	Honest answer: it mattered more to you than to them, and you decided to own it rather than wait. That is self-motivation with a real explanation attached.
How you ensured deadlines were met	A concrete mechanism — reducing decisions to a yes/no, viewing remotely, absorbing the coordination instead of enforcing fairness.
A quantifiable result	Six people housed, before the market cleared out.

> **THE REFLECTION THAT LIFTS IT**
>
> *"Trying to split the work evenly would have been fairer, but it would have failed — because the bottleneck was never effort, it was response time. Absorbing the coordination and reducing everything to a yes-or-no was the only version that actually got six people into a house."*
>
> That is the pragmatic trade-off your own framework asks for: admit it was not the textbook answer, defend why it was right, anchor to the outcome.

If you would rather not use a non-work story here, the alternative is the Pictet project for Q5 and a different, smaller AI example for Q4. I would not do that — Q4 explicitly asks about supporting a teammate or solving a team challenge, and the DDQ story is almost purpose-built for it. Keep the allocation as it is.

20 What never to say in a scenario answer

Don't say	Why
"There's nothing I could really do"	Freezing. There is always something, and the whole question is finding it.
"I'd just fix it myself"	Acting unilaterally on something that is not yours. This is the failure mode firms fear most in juniors.
"I'd wait for my manager"	Passive. Waiting is not a plan when the meeting is today.
"I would never do X"	Absolutist and naive. Real judgment has conditions.
"I'd escalate immediately" — as the first and only move	No judgment about materiality. Escalation after you have established something is different from escalation instead of thinking.
Anything that sounds pleased about being the honest one	Sanctimony reads terribly on camera. Be practical, not righteous.
Running long to show you have thought of everything	These are pass-or-fail. Hit the four beats and stop talking.

21 Before you record

1. **Say the four beats out loud once**, in order, against a scenario you make up on the spot. Not to get it right — to prove to yourself the structure holds on something unfamiliar. That is the actual cure for the panic.
2. **For Q2, lead with the stake.** "Late is an inconvenience, wrong is a trust problem." If you only remember one sentence from Part IV, make it that one.
3. **For Q3, split the sentence into three.** Tool, review, disclosure. Then give the alternatives, because the question asks for them.
4. **Keep both scenario answers under 90 seconds.** They are shorter than your story answers by design.
5. **Check the allocation holds:** Pictet to Q4, house hunt to Q5. Do not let the DDQ story leak into both.

A note on what this document is. The six prompts are as you supplied them. Everything else here — the four-beat protocol (carried over from doc 46), the five scoring criteria, the control/no-control split, the three traps in question 2, the three-way separation in question 3, the four alternatives, the story allocation, and both scripted answers — is my analysis rather than sourced fact, and should be treated as a way of structuring your thinking rather than as information about Goldman's actual scoring. The one factual claim carried in from elsewhere is doc 54's point about plausible-but-wrong output, which stands on its own. Do not assert anything about how Goldman marks these; nobody outside the firm knows, and the structure works regardless.